

Journey to achieving Cloud Native at Kakao

정기훈(lev.jung@kakaocorp.com)
클라우드네이티브파트, 카카오

발표자 소개

정기훈

#SoftwareEngineer#Cloud#LargeScale#Networking#Automation#CICD#SRE

현) 클라우드네이티브파트 파트장, 카카오
전) SW R&D Center, 삼성전자

01 What is Cloud Native?

02 Introducing Kakao Cloud

03 How Kakao achieved to Cloud Native

04 Summary

01 What is Cloud Native?

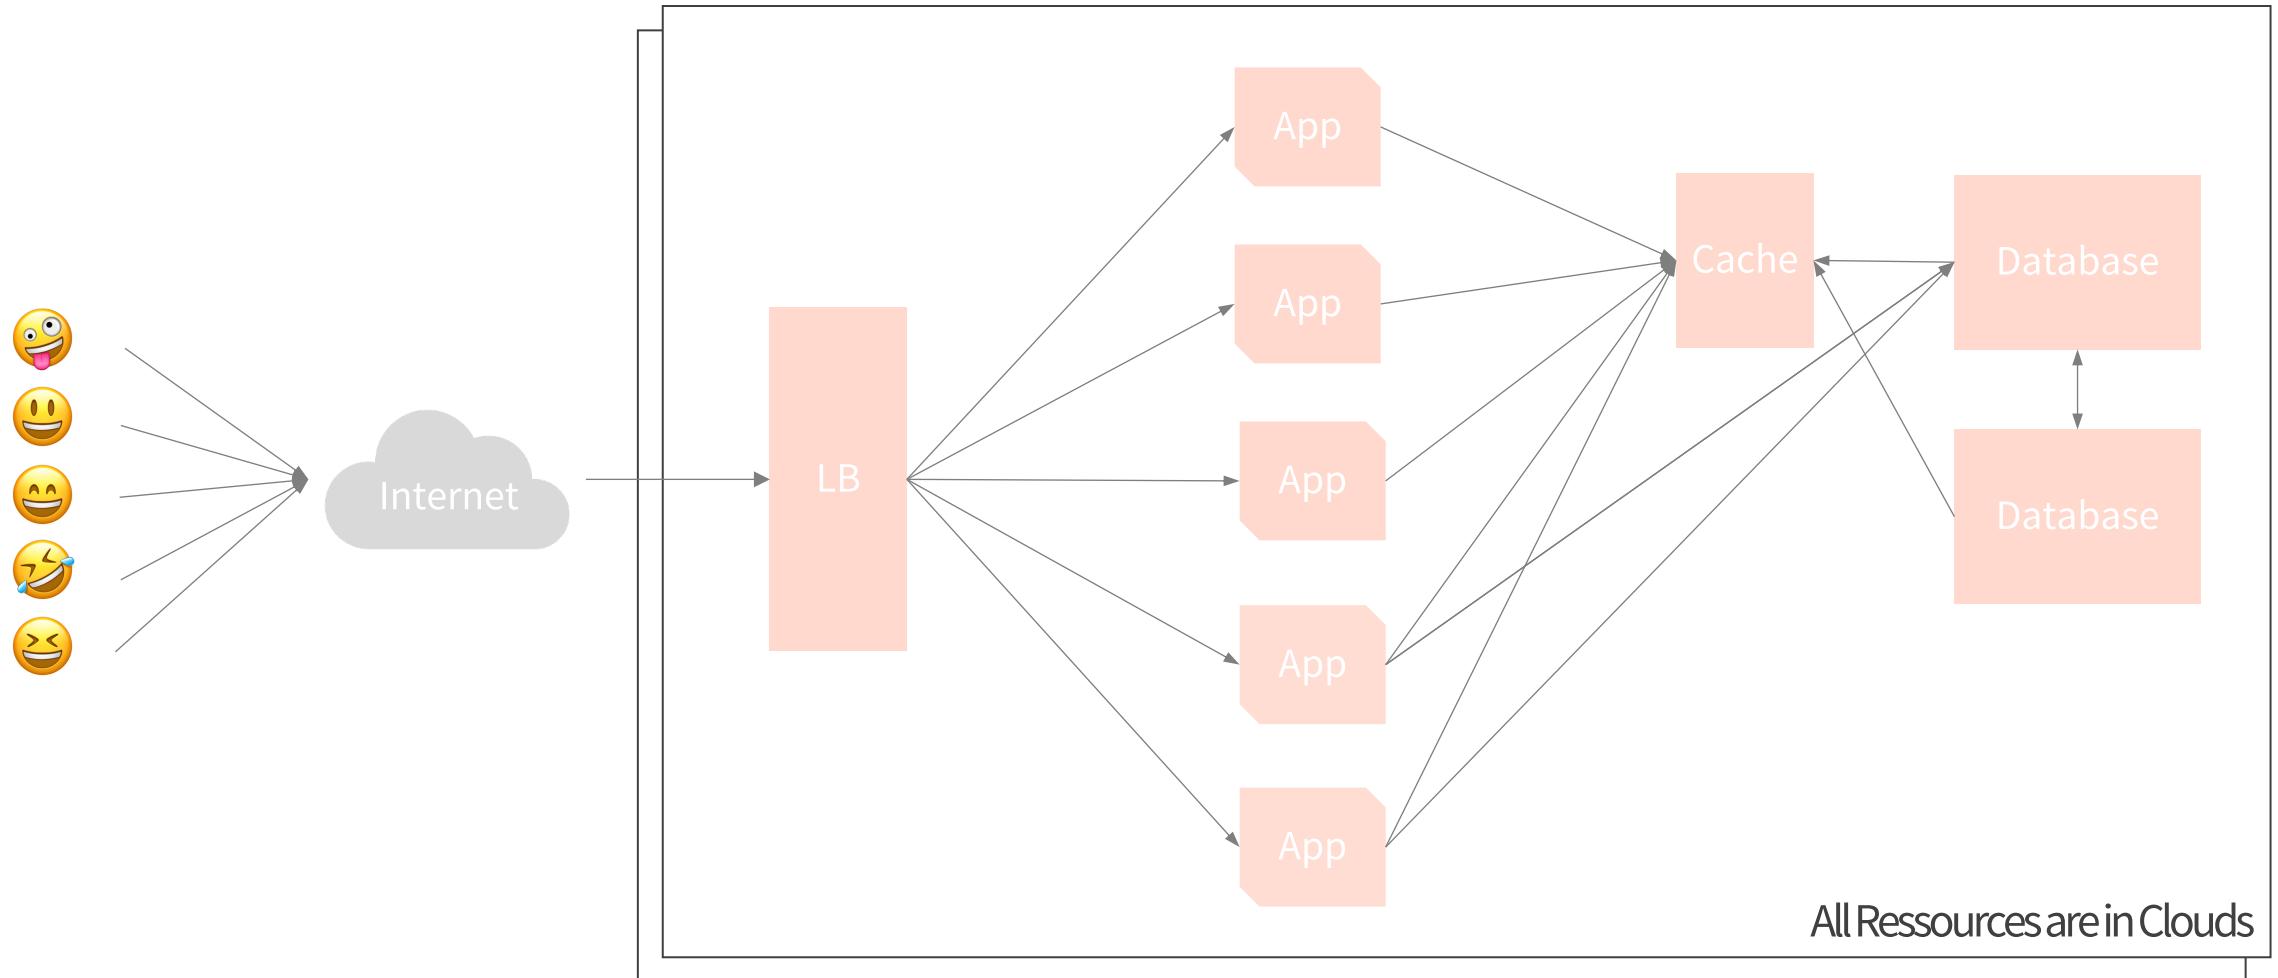
Cloud Native Definition from CNCF

Empower organizations to build and run **scalable applications in clouds.**

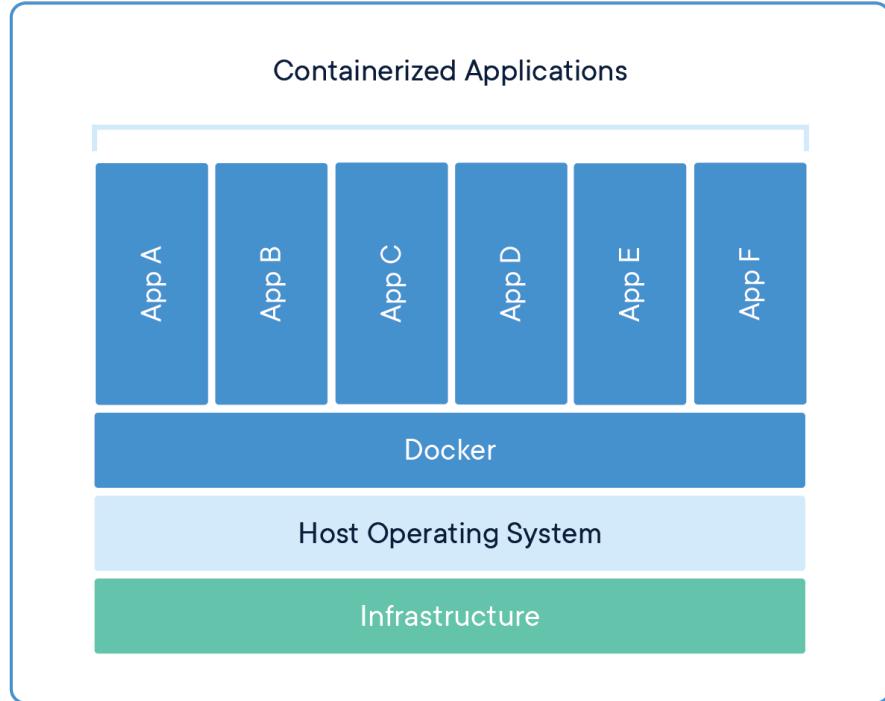
Exemplify approach:

- containers
- microservices
- service meshes
- immutable infrastructure
- declarative APIs

Scalable Application in clouds

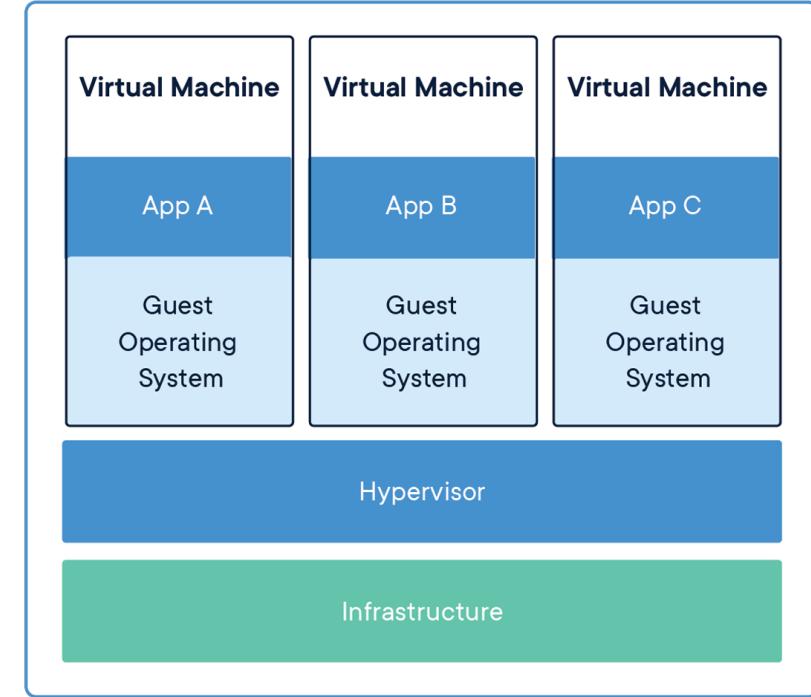


Container



Container

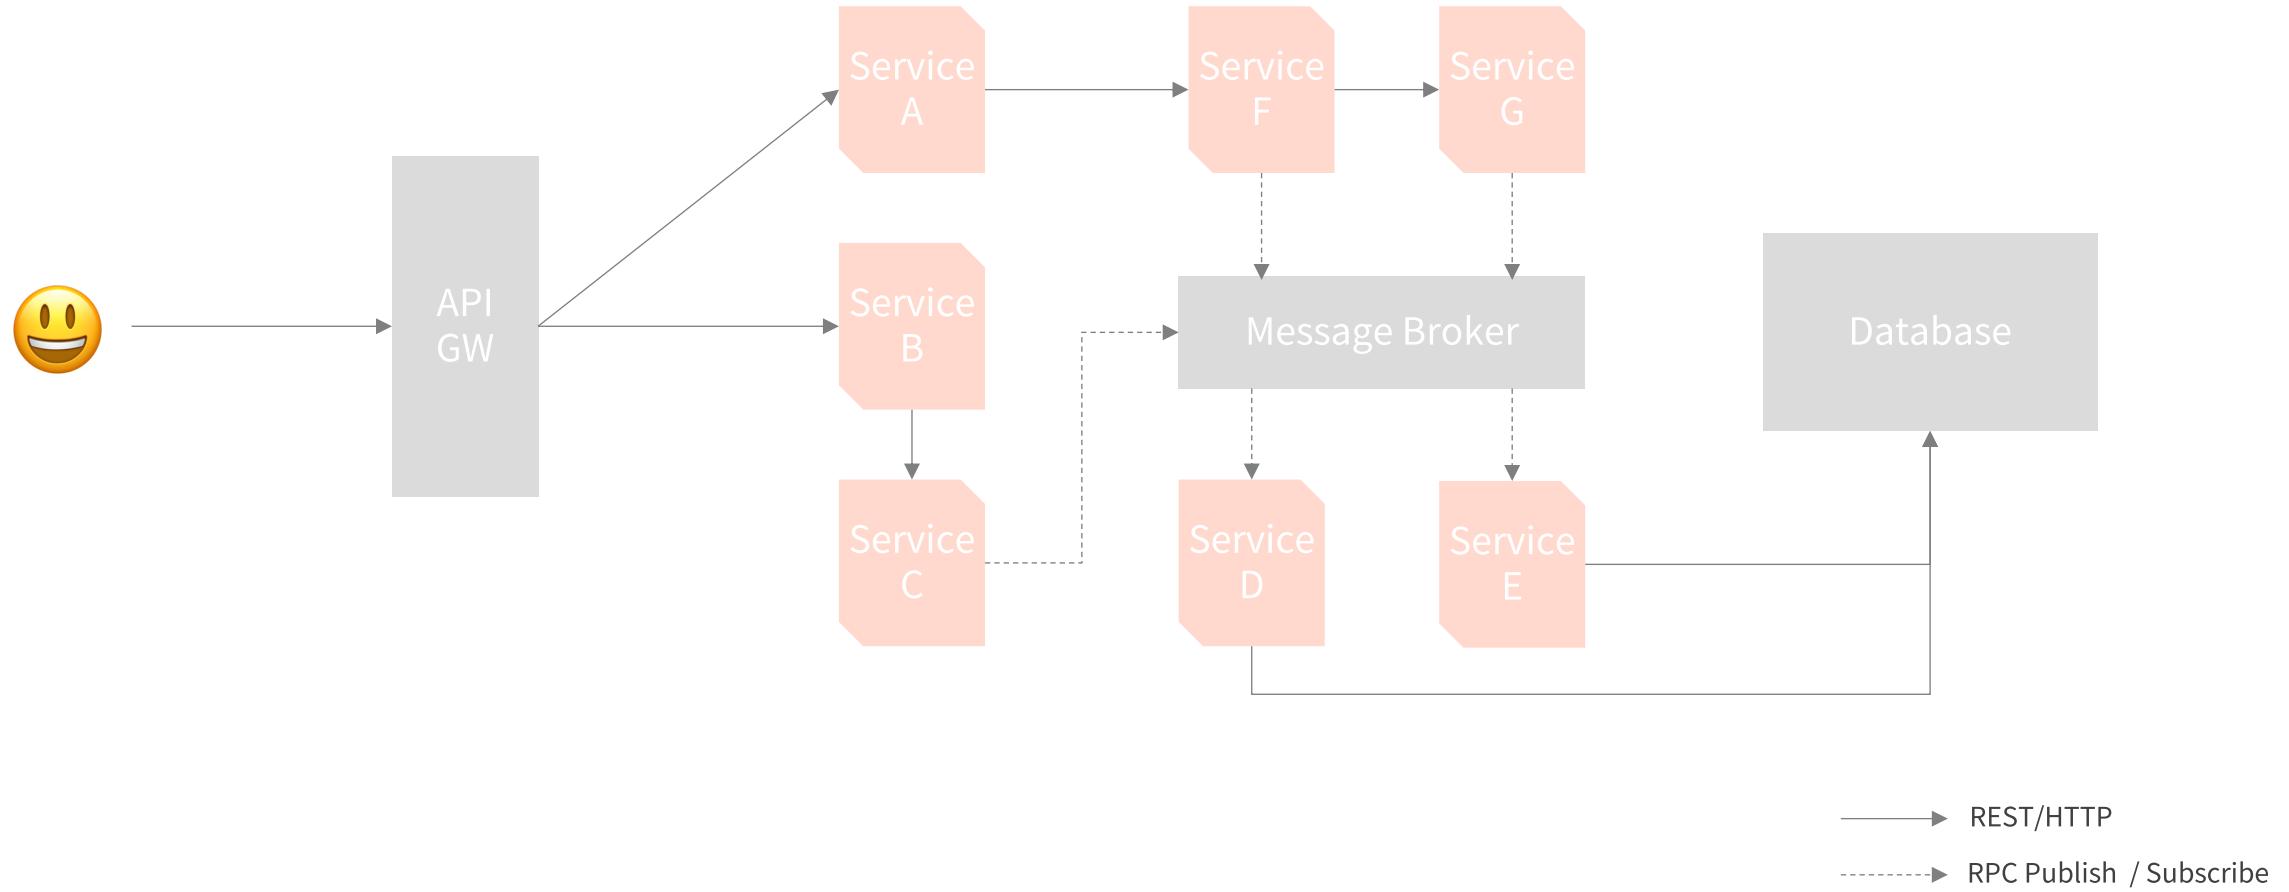
Host Kernel을 공유
User Space 격리 제공
레이어 기반의 이미지 구성



Virtual Machine

Host Kernel과 완전 격리 가능
Guest OS의 선택이 가능
디스크 기반의 이미지 구성

Microservice Architecture

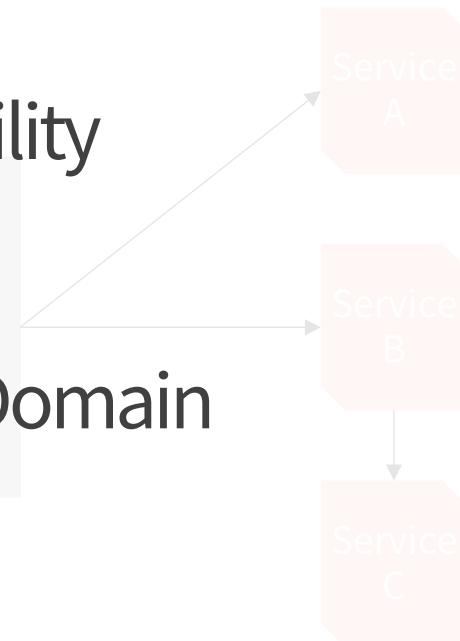


Microservice Architecture

Pros

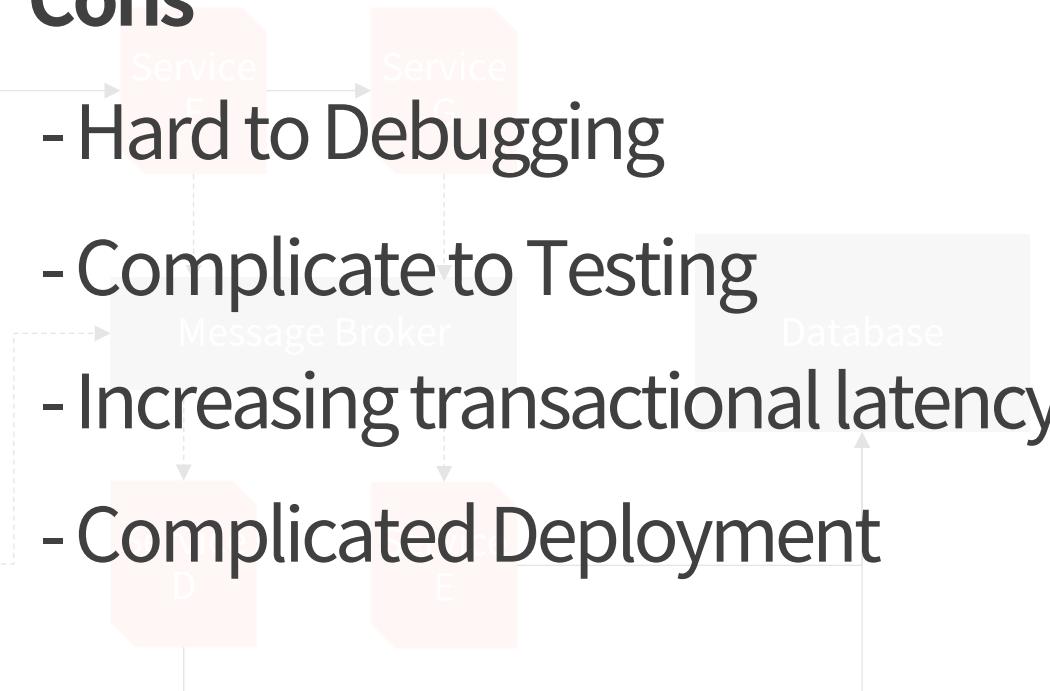
- Single Responsibility
- Easy to Scale

- Isolating Failure Domain
- Agility to Deploy



Cons

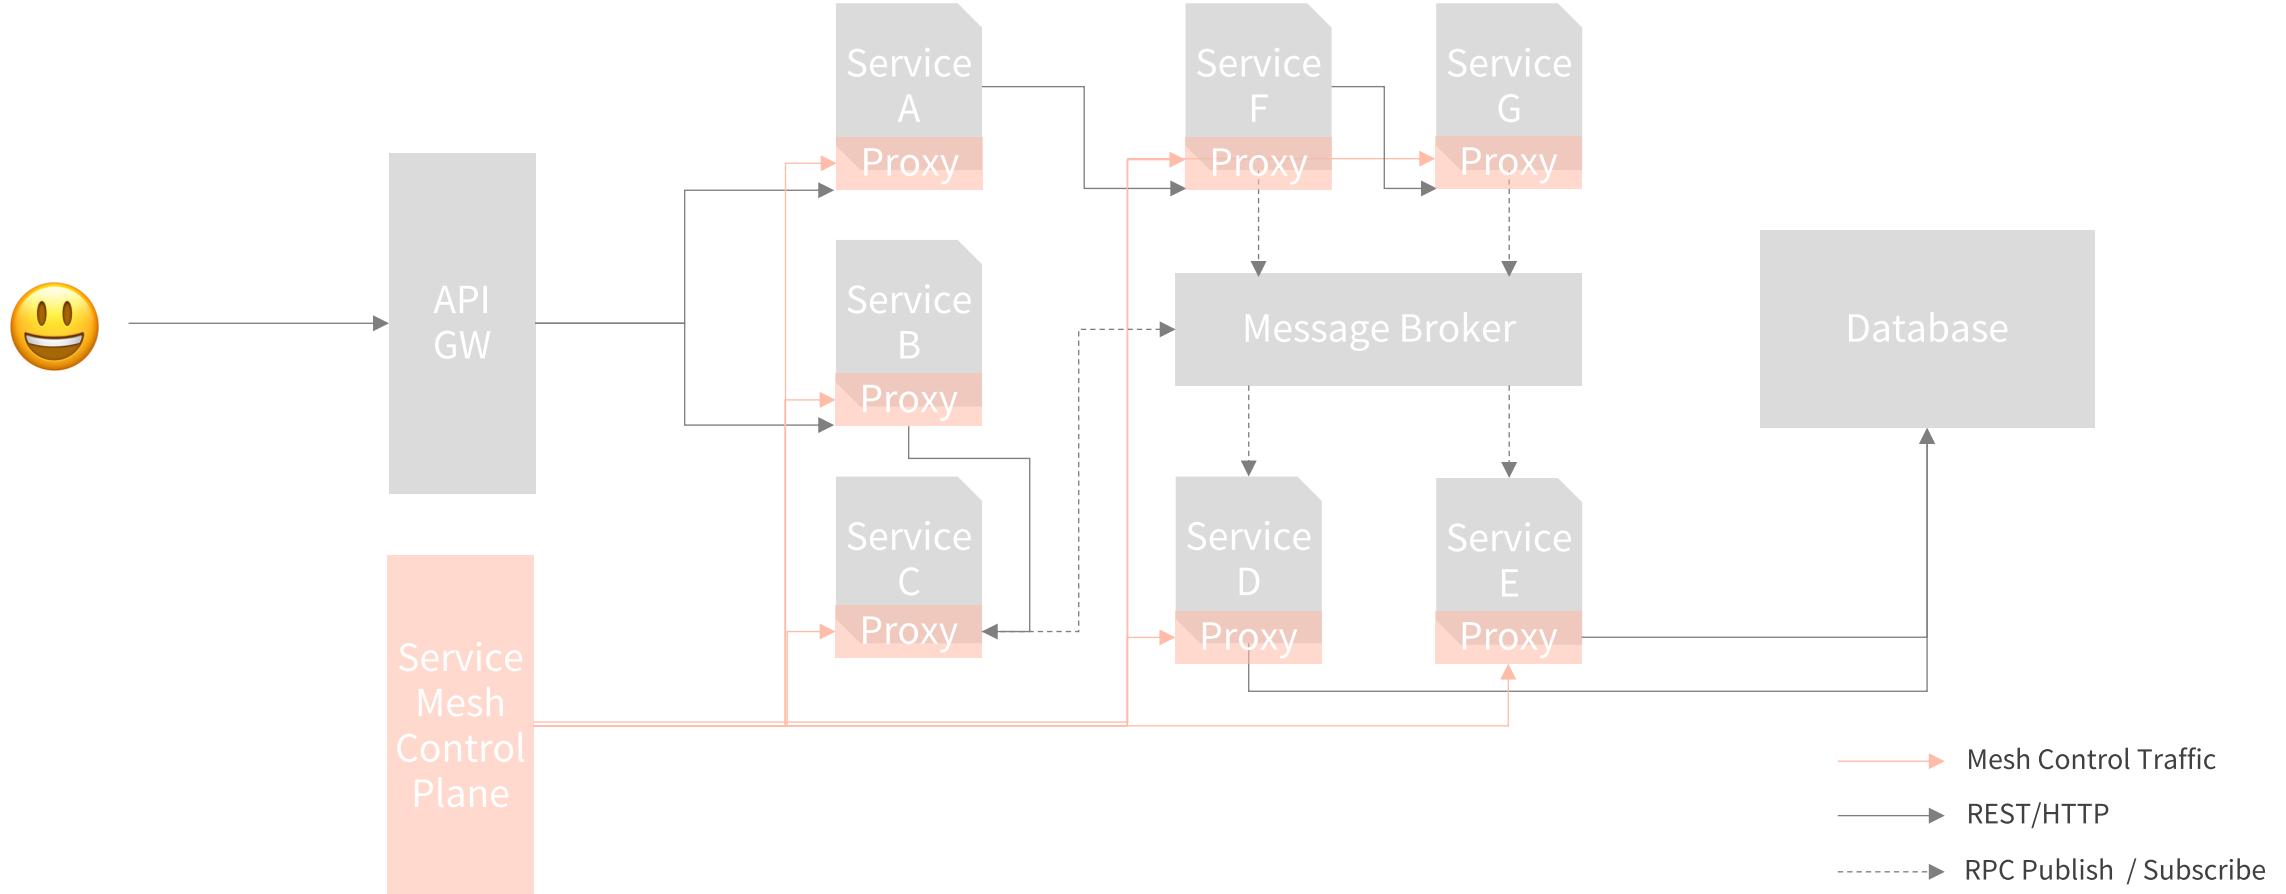
- Hard to Debugging
- Complicate to Testing
- Increasing transactional latency
- Complicated Deployment



→ REST/HTTP

→ RPC Publish / Subscribe

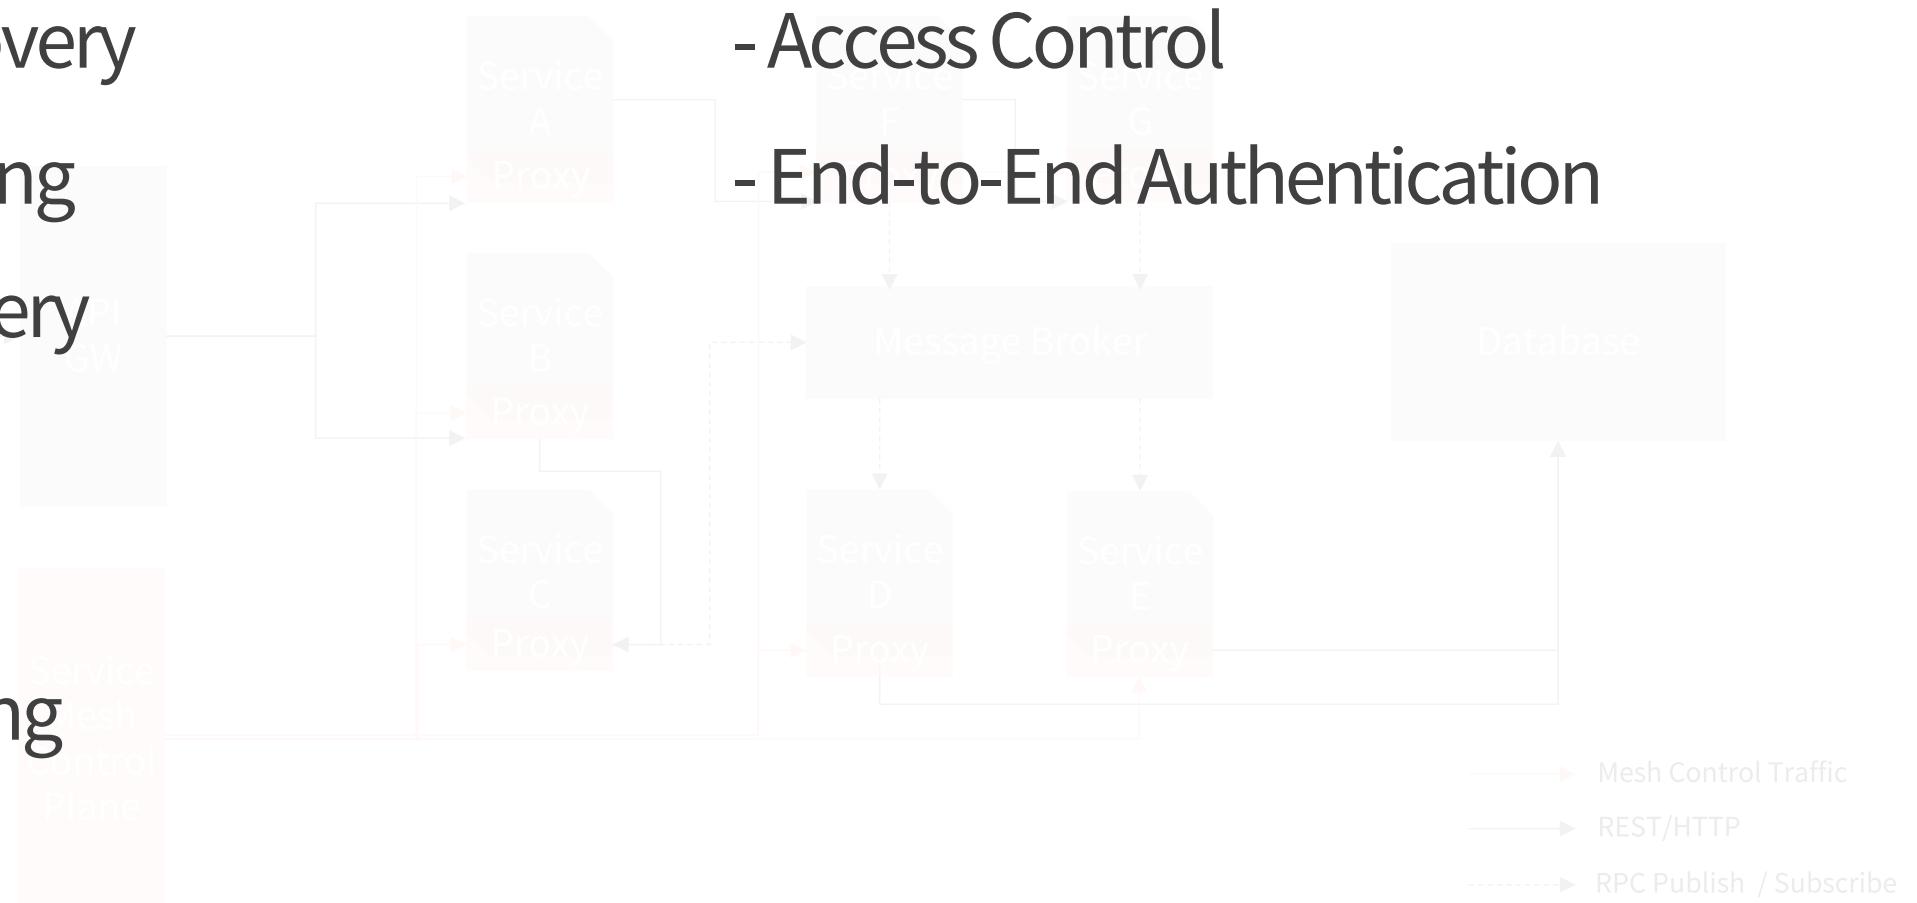
Service Mesh



Service Mesh

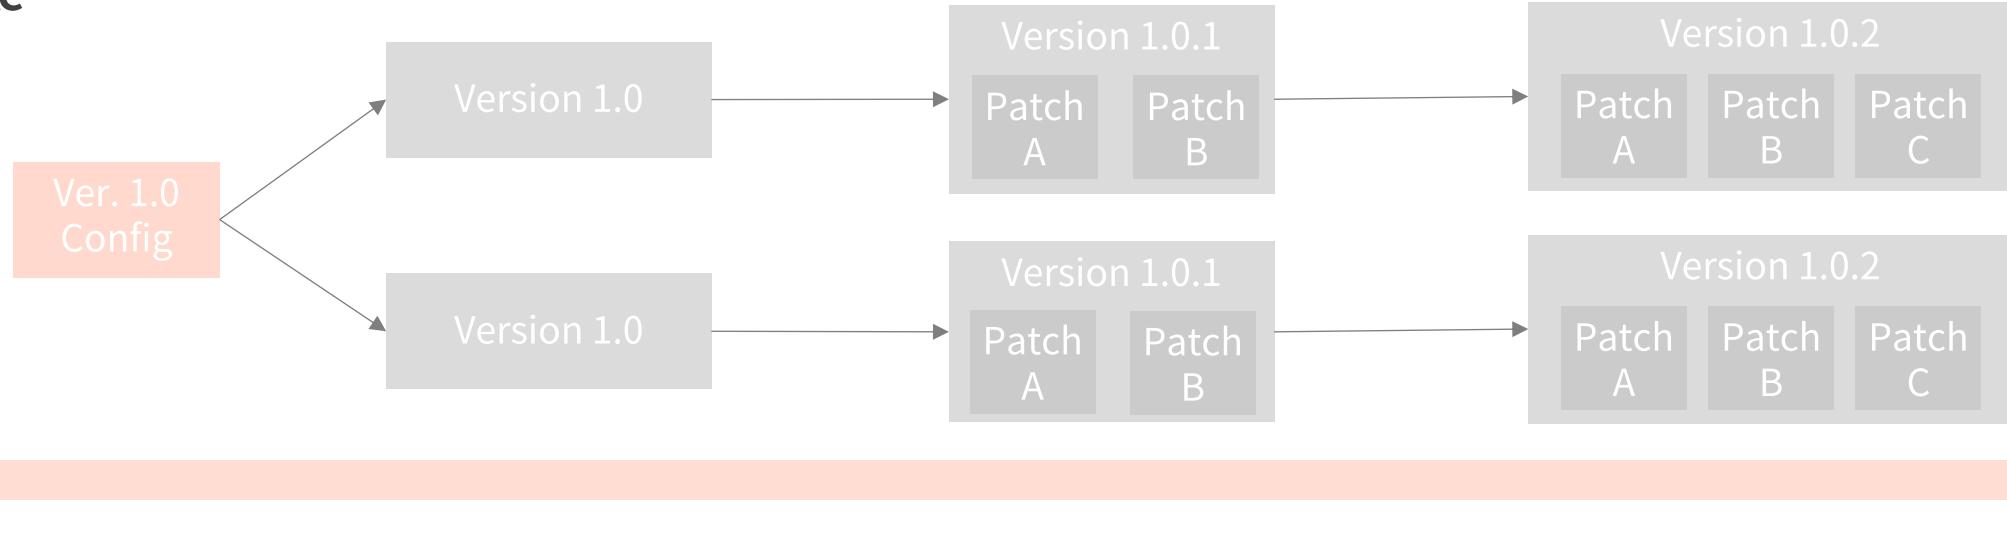
Features

- Service Discovery
- Load Balancing
- Failure recovery
- Metric
- Monitoring
- Traffic Steering

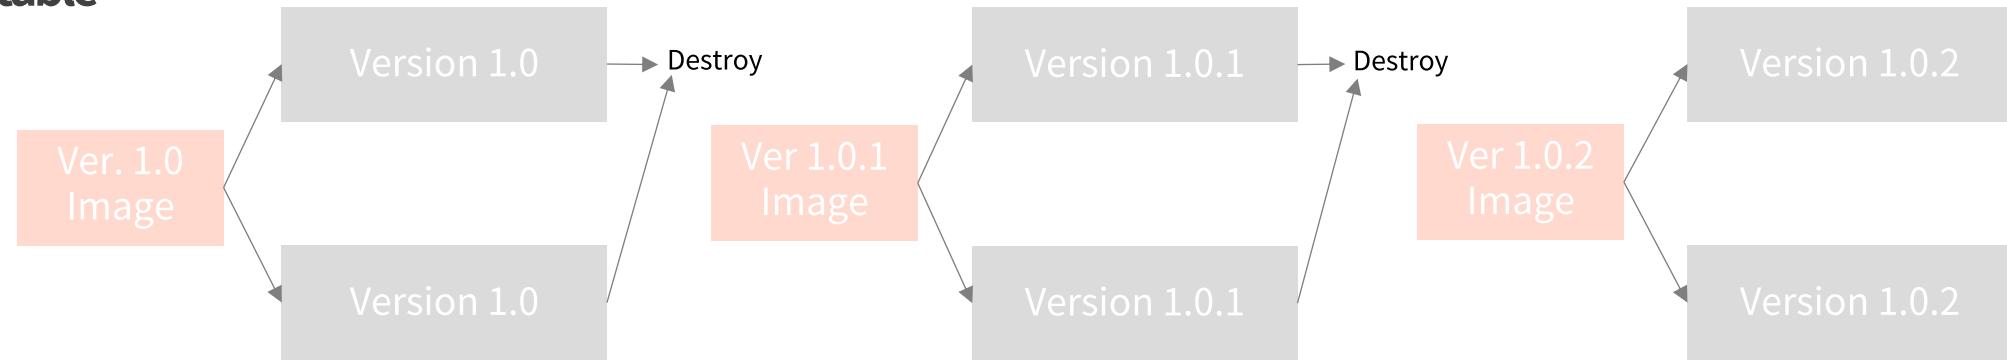


Immutable Infrastructure

Mutable



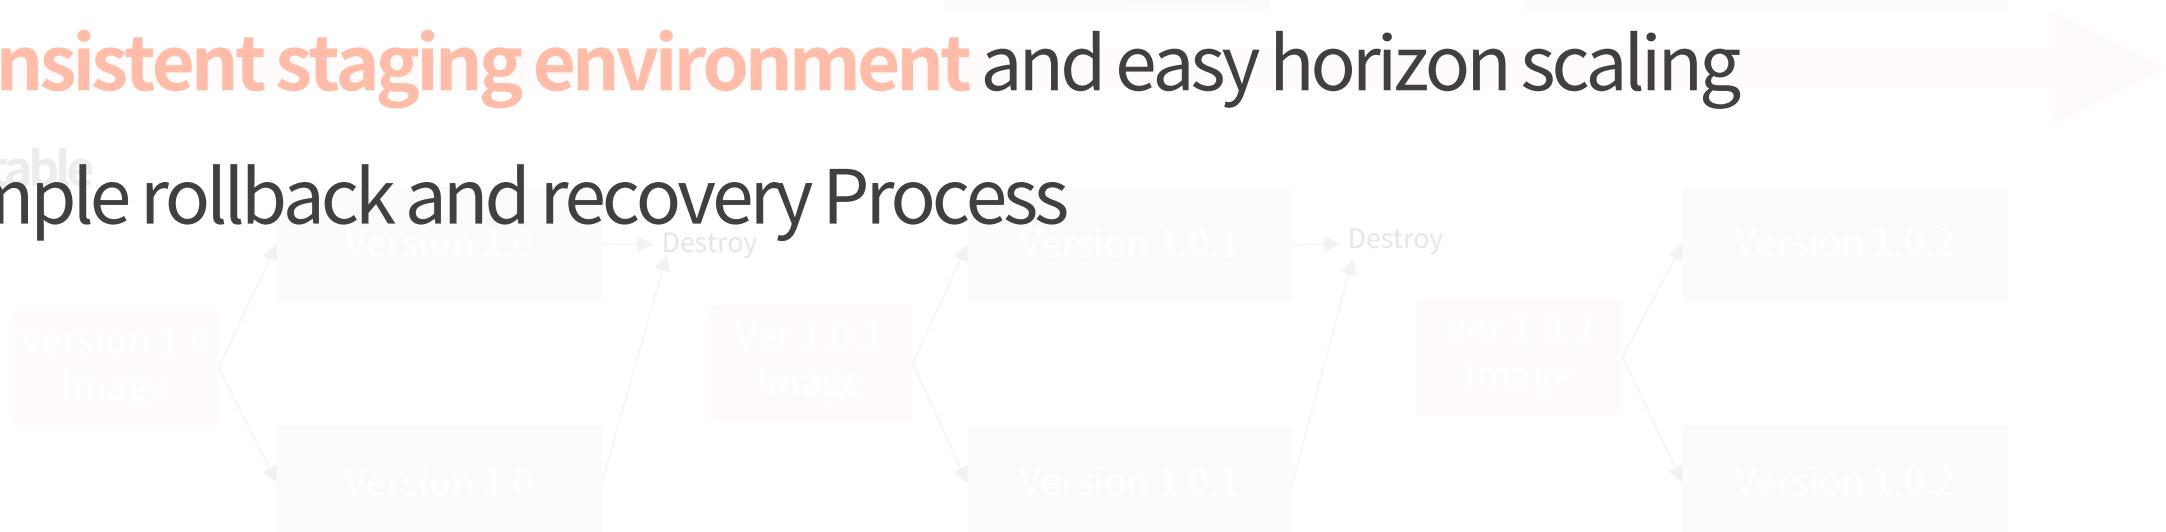
Immutable



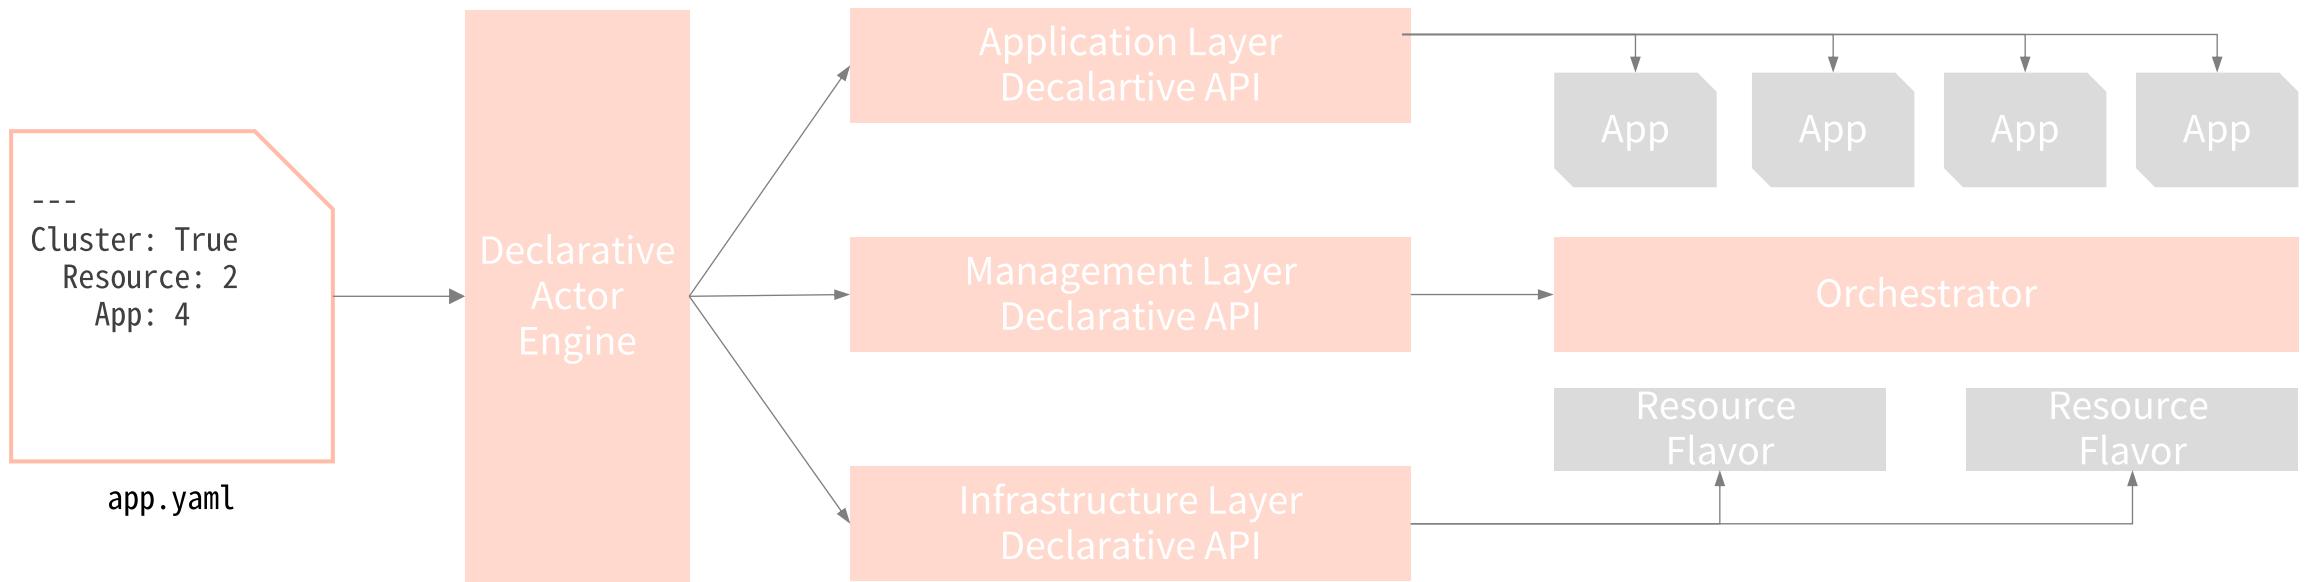
Immutable Infrastructure

Advantage

- **No Configuration Drift** or Snowflake servers
- Known-good server state and fewer deployment failure
- **Consistent staging environment** and easy horizon scaling
- Simple rollback and recovery Process



Declarative API



Declarative API

Characteristic

- Defining Object as an Application or Infrastructure
- Defining desire state of Object
- CRUD(Create, Read, Update, Delete) on Object
- Document is self-explainable and Human readable

Additionally, Twelve Factors

1. Codebase

One codebase tracked in revision control

2. Dependencies

Explicitly declare and isolate dependencies

3. Config

Store config in the environment

4. Backing Service

Treat backing services as attached resources

5. Build, Release, Run

Strictly separate build and run stages

6. Process

Execute the app as one or more stateless processes

7. Port Biding

Export services via port binding

8. Concurrency

Scale out via the process model

9. Disposability

Maximize robustness with fast startup and graceful shutdown

10. Dev/prod parity

Keep development, staging, and production as similar as possible

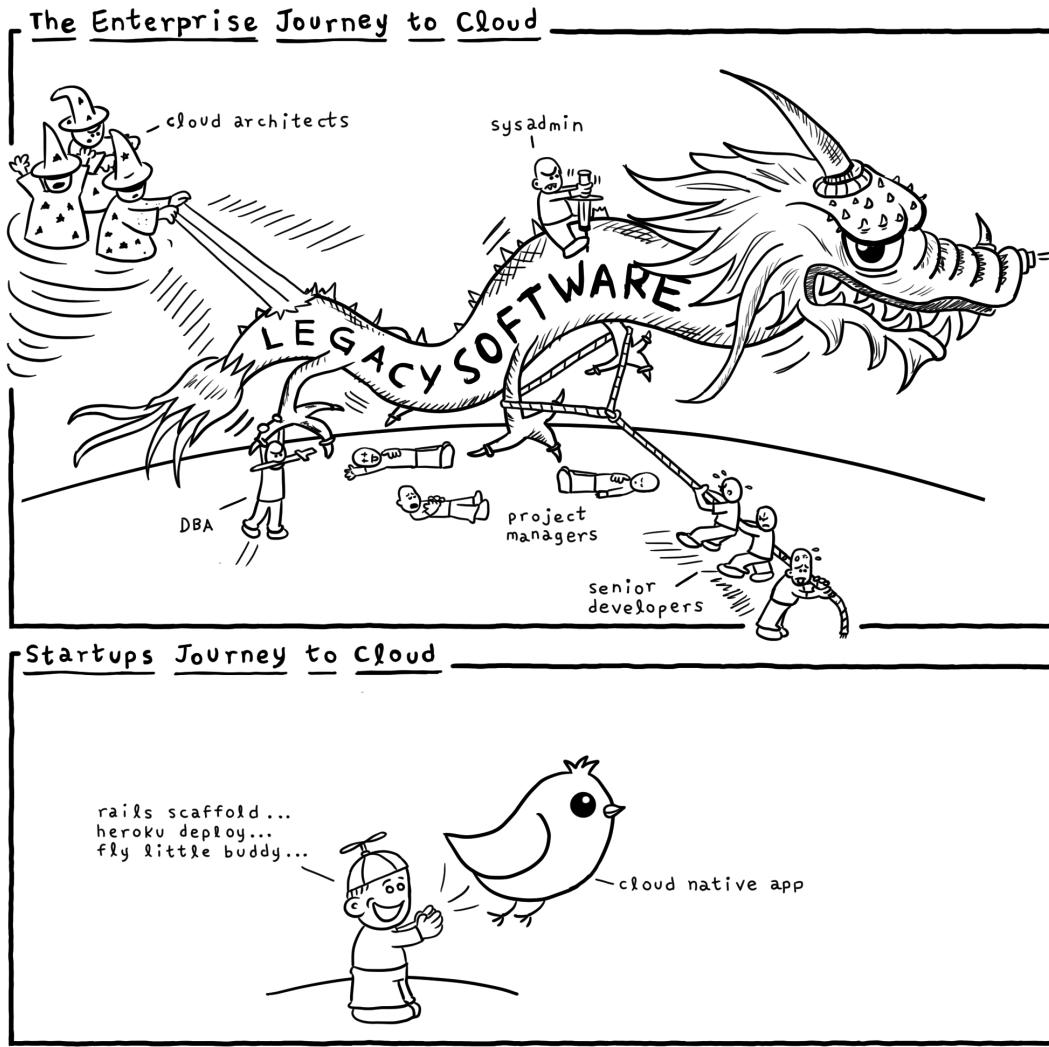
11. Log

Treat logs as event streams

12. Admin processes

Run admin/management tasks as one-off processes

However,



Daniel Stori {turnoff.us}
Thanks to Michael Tharrington

But People say,

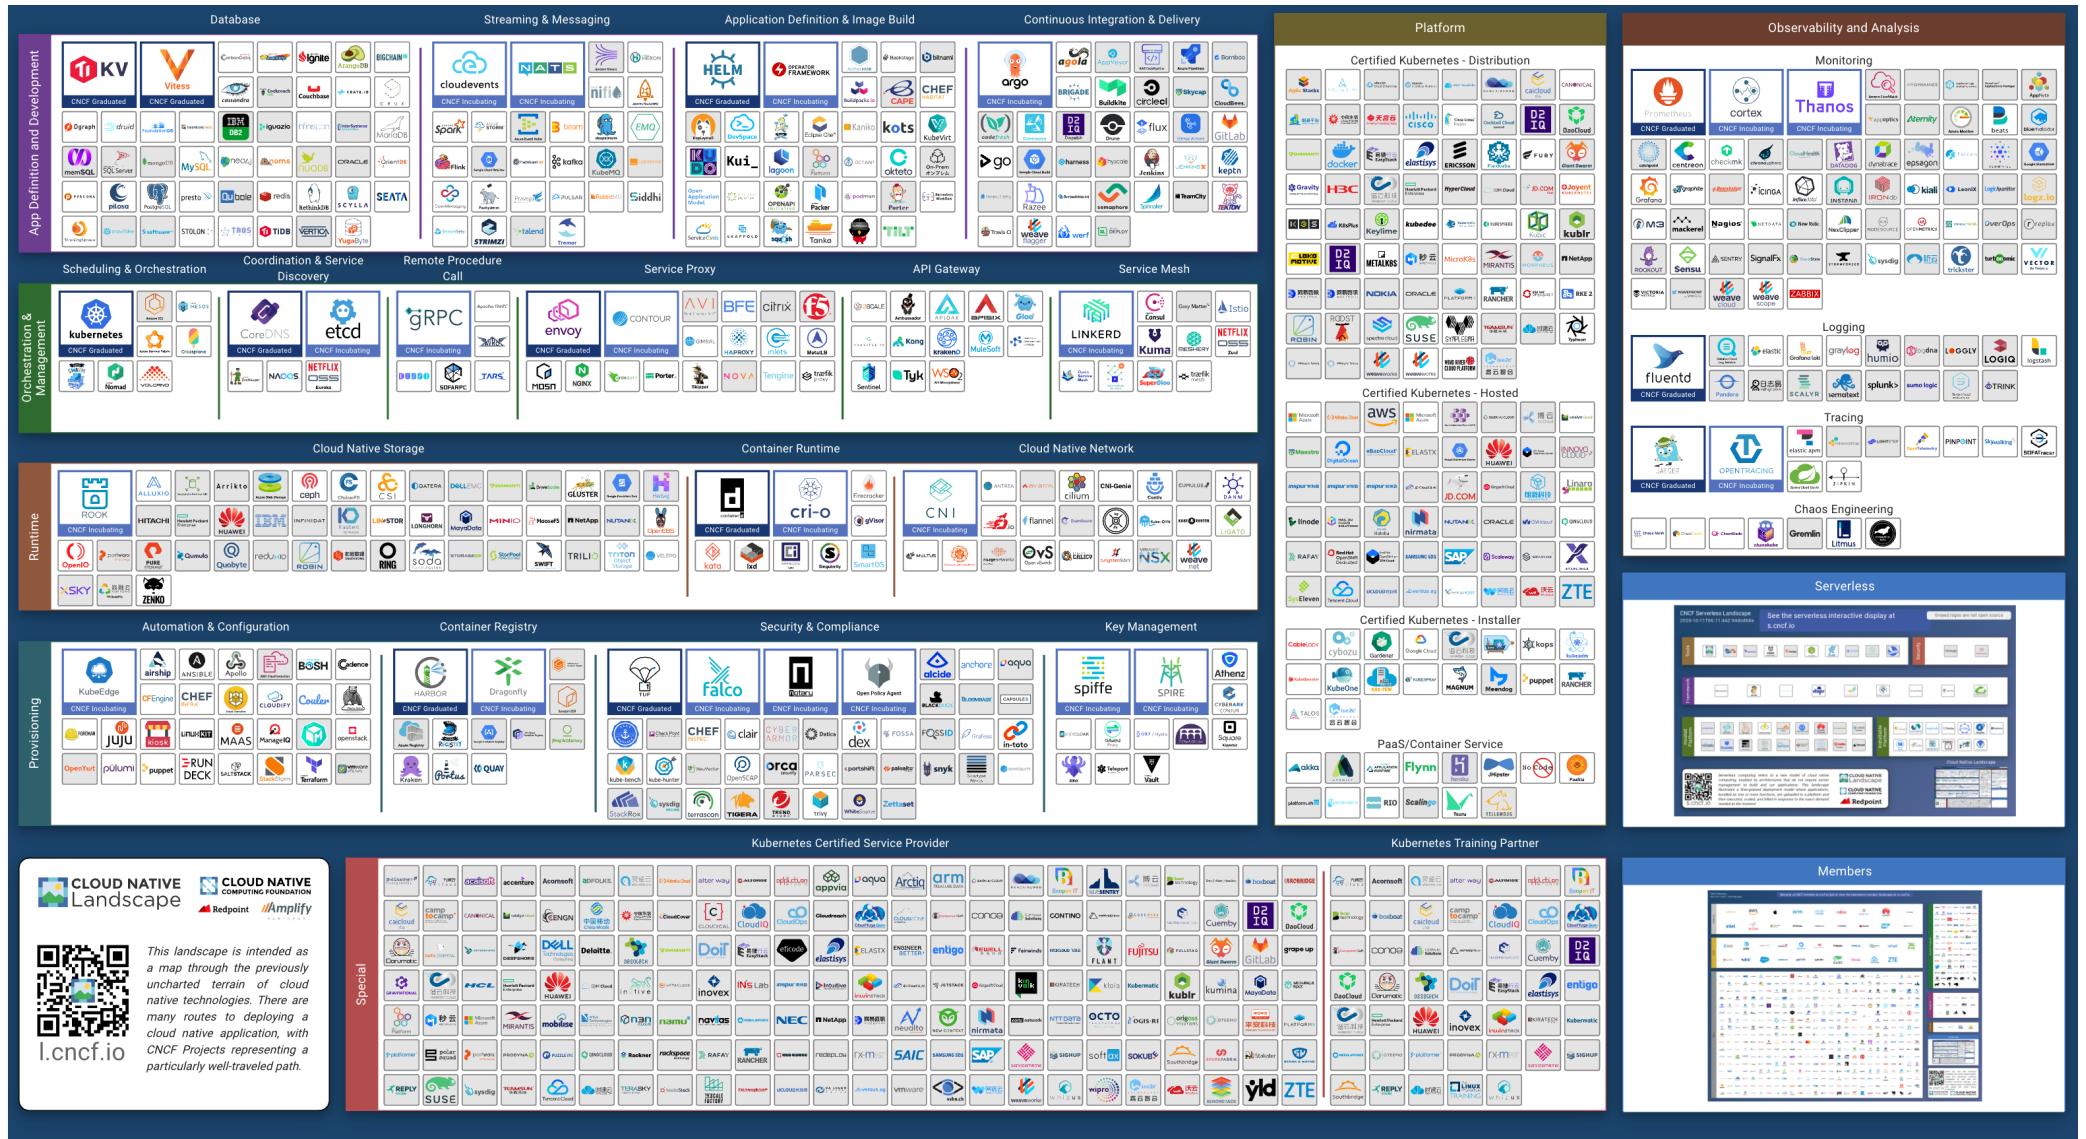
“그래서,

이제는 쿠버네티스만 있으면 다 되는 거죠!?”

...



Kubernetes rules them all?



<https://landscape.cncf.io>

Long Journey

1. CONTAINERIZATION

- Commonly done with Docker containers
- Any size application and dependencies (even PDP-11 code running on an emulator) can be containerized
- Over time, you should aspire towards splitting suitable applications and writing future functionality as microservices

3. ORCHESTRATION & APPLICATION DEFINITION

- Kubernetes is the market-leading orchestration solution
- You should select a Certified Kubernetes Distribution, Hosted Platform, or Installer: cncf.io/ck
- Helm Charts help you define, install, and upgrade even the most complex Kubernetes application

5. SERVICE PROXY, DISCOVERY, & MESH

- CoreDNS is a fast and flexible tool that is useful for service discovery
- Envoy and Linkerd each enable service mesh architectures
- They offer health checking, routing, and load balancing

2. CI/CD

- Setup Continuous Integration/Continuous Delivery (CI/CD) so that changes to your source code automatically result in a new container being built, tested, and deployed to staging and eventually, perhaps, to production
- Setup automated rollouts, roll backs and testing
- Argo is a set of Kubernetes-native tools for deploying and running jobs, applications, workflows, and events using GitOps paradigms such as continuous and progressive delivery and MLOps


Argo
CNCF Incubating

4. OBSERVABILITY & ANALYSIS

- Pick solutions for monitoring, logging and tracing
- Consider CNCF projects Prometheus for monitoring, Fluentd for logging and Jaeger for Tracing
- For tracing, look for an OpenTracing-compatible implementation like Jaeger


Prometheus
CNCF Graduated


fluentd
CNCF Graduated


JAAGER
CNCF Graduated


OPENTRACING
CNCF Incubating

A winding path through a landscape with trees, a lake, and a cliff, serving as a metaphor for the 'Long Journey' of cloud native development.

- Envoy and Linkerd each enable service mesh architectures
- They offer health checking, routing, and load balancing



CNCF Graduated



CNCF Graduated



CNCF Incubating



7. DISTRIBUTED DATABASE & STORAGE

When you need more resiliency and scalability than you can get from a single database, Vitess is a good option for running MySQL at scale through sharding. Rook is a storage orchestrator that integrates a diverse set of storage solutions into Kubernetes. Serving as the "brain" of Kubernetes, etcd provides a reliable way to store data across a cluster of machines. TiKV is a high performant distributed transactional key-value store written in Rust.



CNCF Graduated



CNCF Incubating



CNCF Incubating



CNCF Incubating



9. CONTAINER REGISTRY & RUNTIME

Harbor is a registry that stores, signs, and scans content. You can use alternative container runtimes. The most common, both of which are OCI-compliant, are containerd and CRI-O.



CNCF Graduated



CNCF Incubating



CNCF Incubating



JAEGER



OPENTRACING

CNCF Graduated

CNCF Incubating



CLOUD NATIVE

6. NETWORKING, POLICY, & SECURITY

To enable more flexible networking, use a CNI-compliant network project like Calico, Flannel, or Weave Net. Open Policy Agent (OPA) is a general-purpose policy engine with uses ranging from authorization and admission control to data filtering. Falco is an anomaly detection engine for cloud native.



CNCF Incubating



CNCF Incubating



CNCF Incubating

8. STREAMING & MESSAGING

When you need higher performance than JSON-REST, consider using gRPC or NATS. gRPC is a universal RPC framework. NATS is a multi-modal messaging system that includes request/reply, pub/sub and load balanced queues. CloudEvents is a specification for describing event data in common ways.



CNCF Incubating



CNCF Incubating



CNCF Incubating

10. SOFTWARE DISTRIBUTION

If you need to do secure software distribution, evaluate Notary, an implementation of The Update Framework.



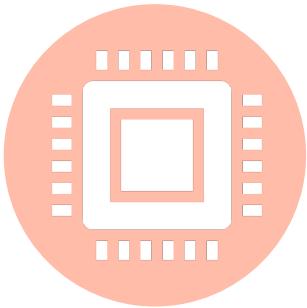
CNCF Graduated



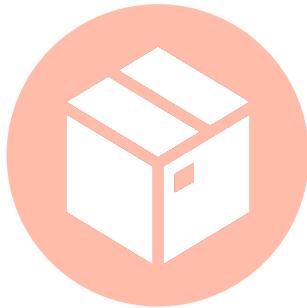
CNCF Incubating

02 Introducing Kakao Cloud

Some Numbers



PM
0000



VM
0000



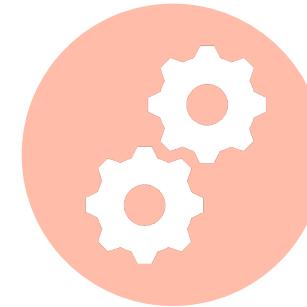
VOLUME
 \sim 0000



PROJECT
0000

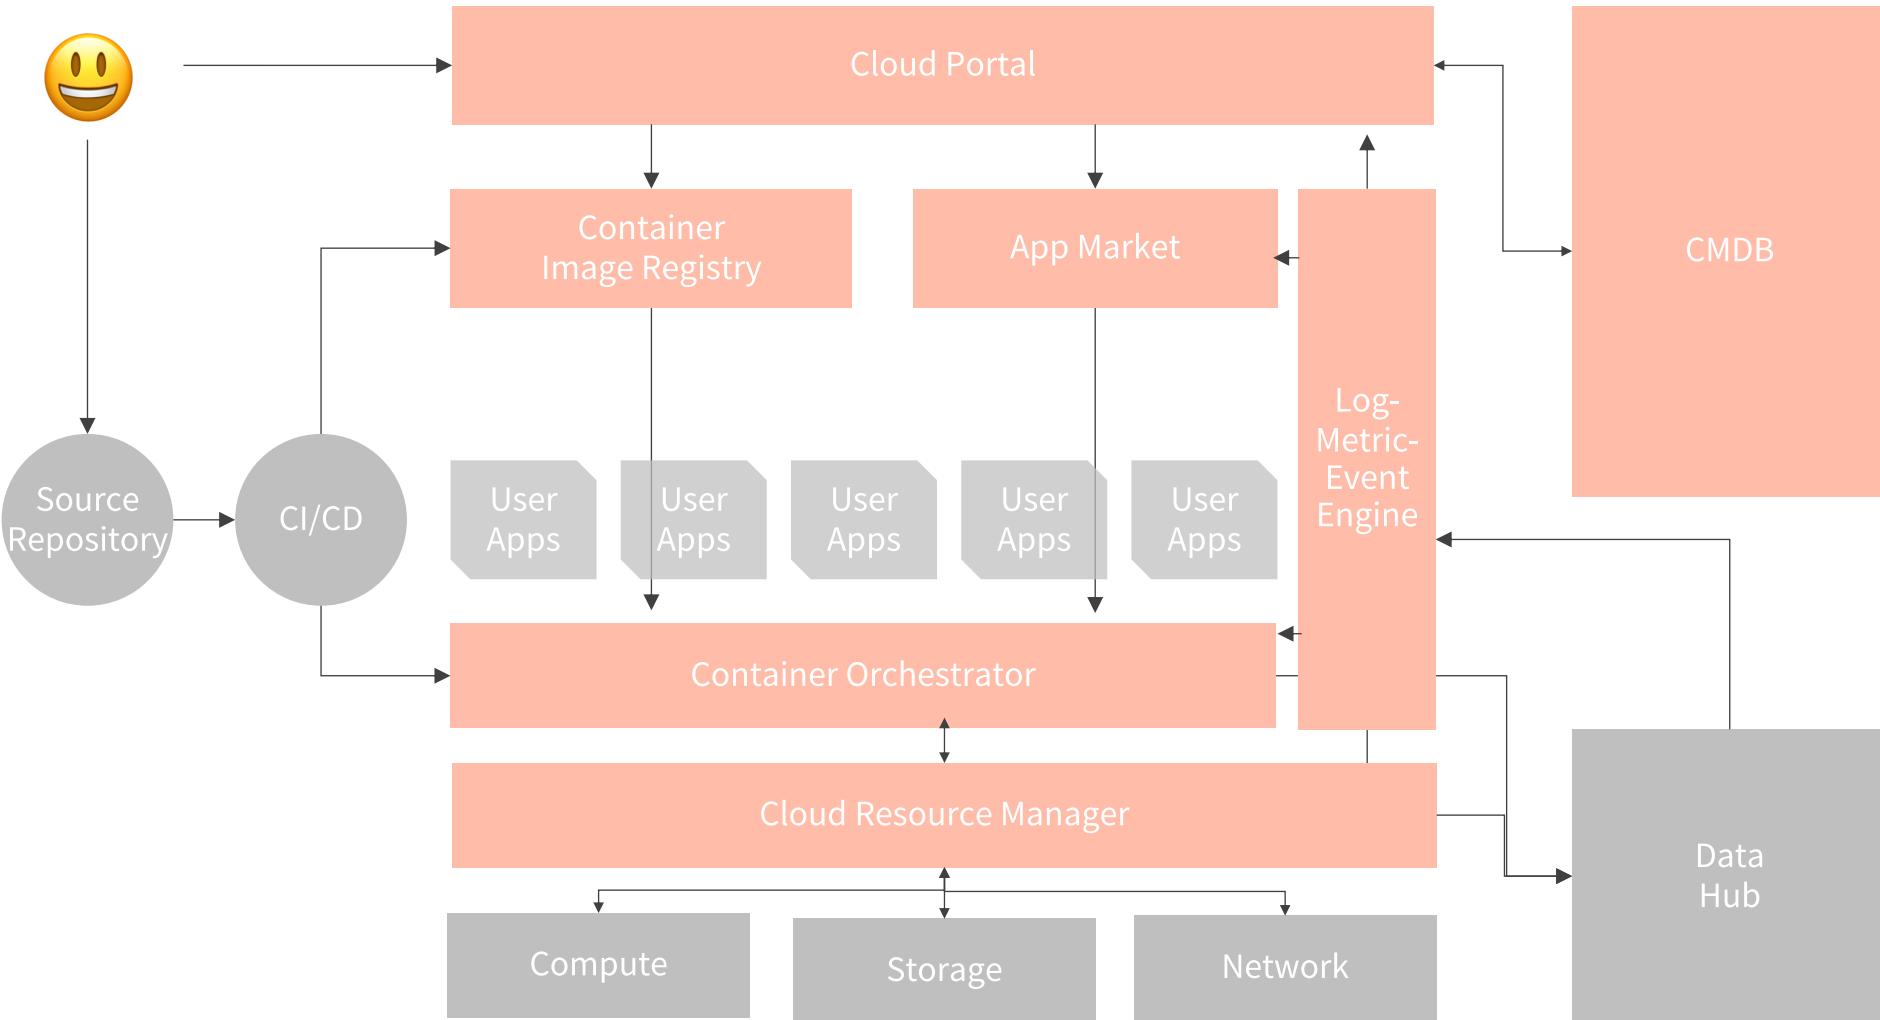


USER
0000

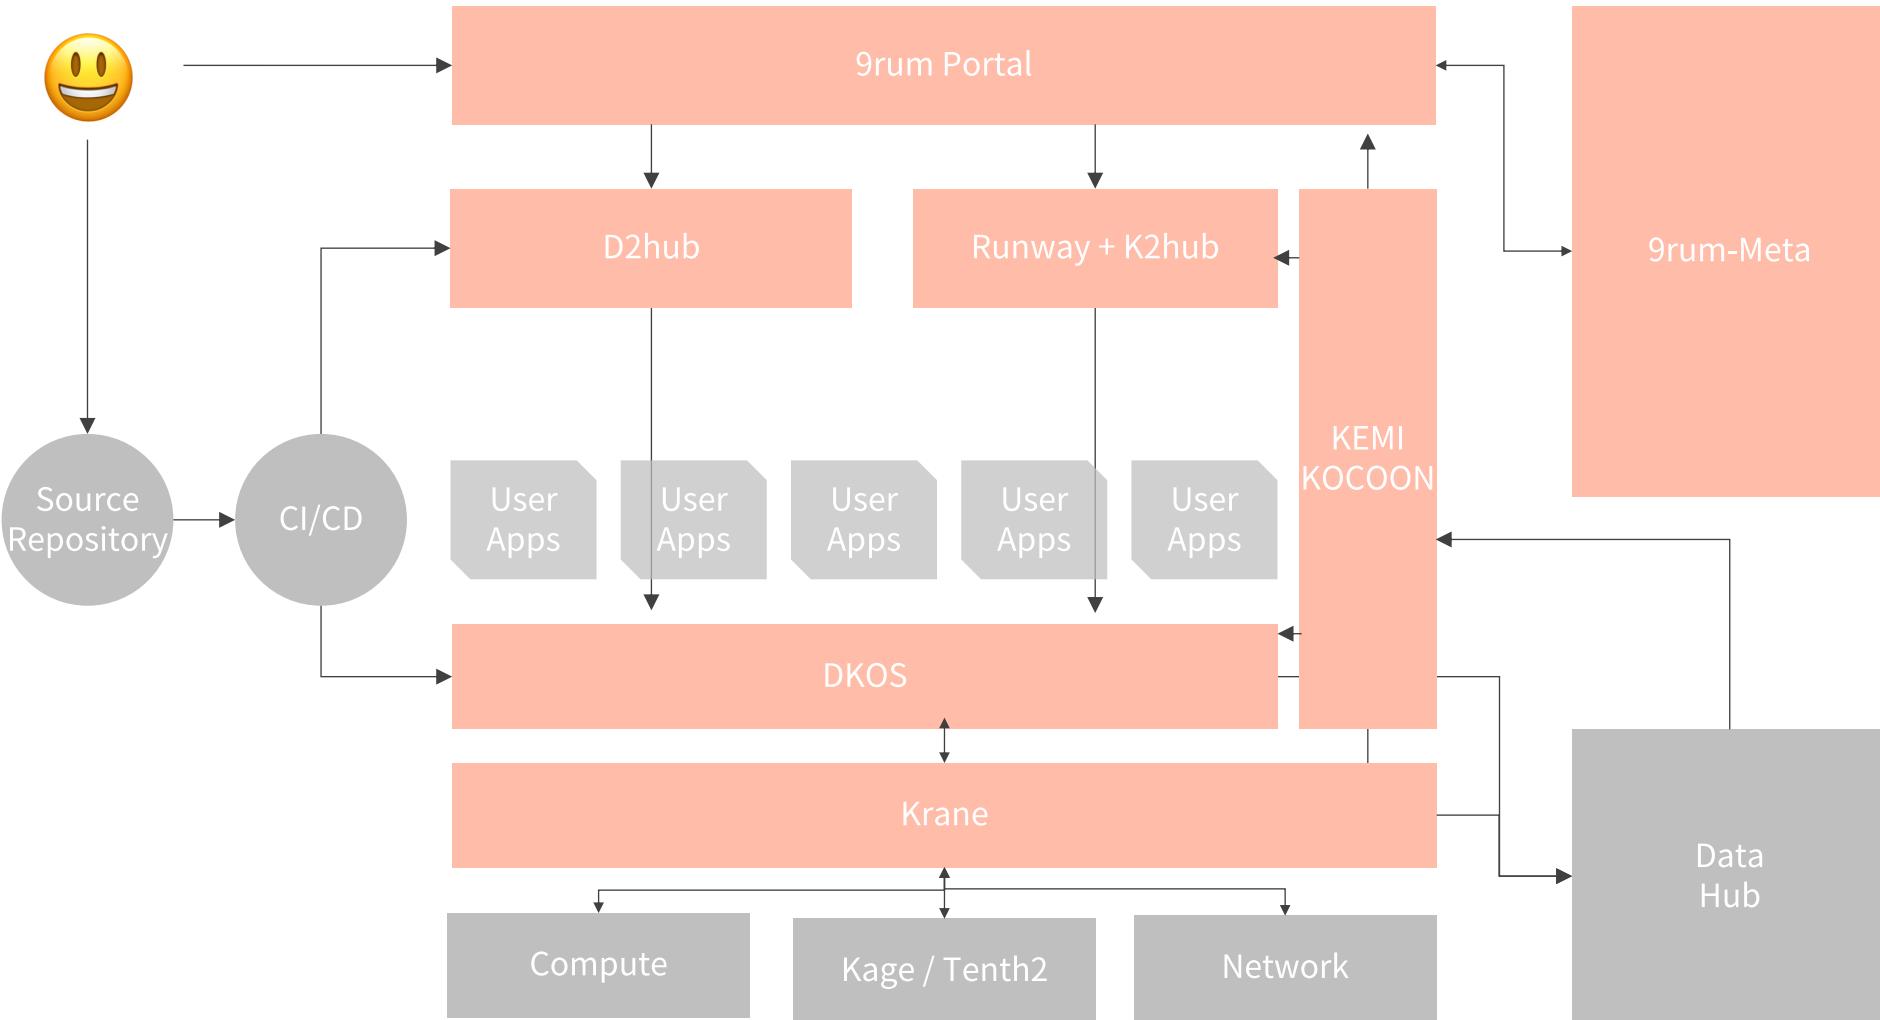


PRODUCTION
0000

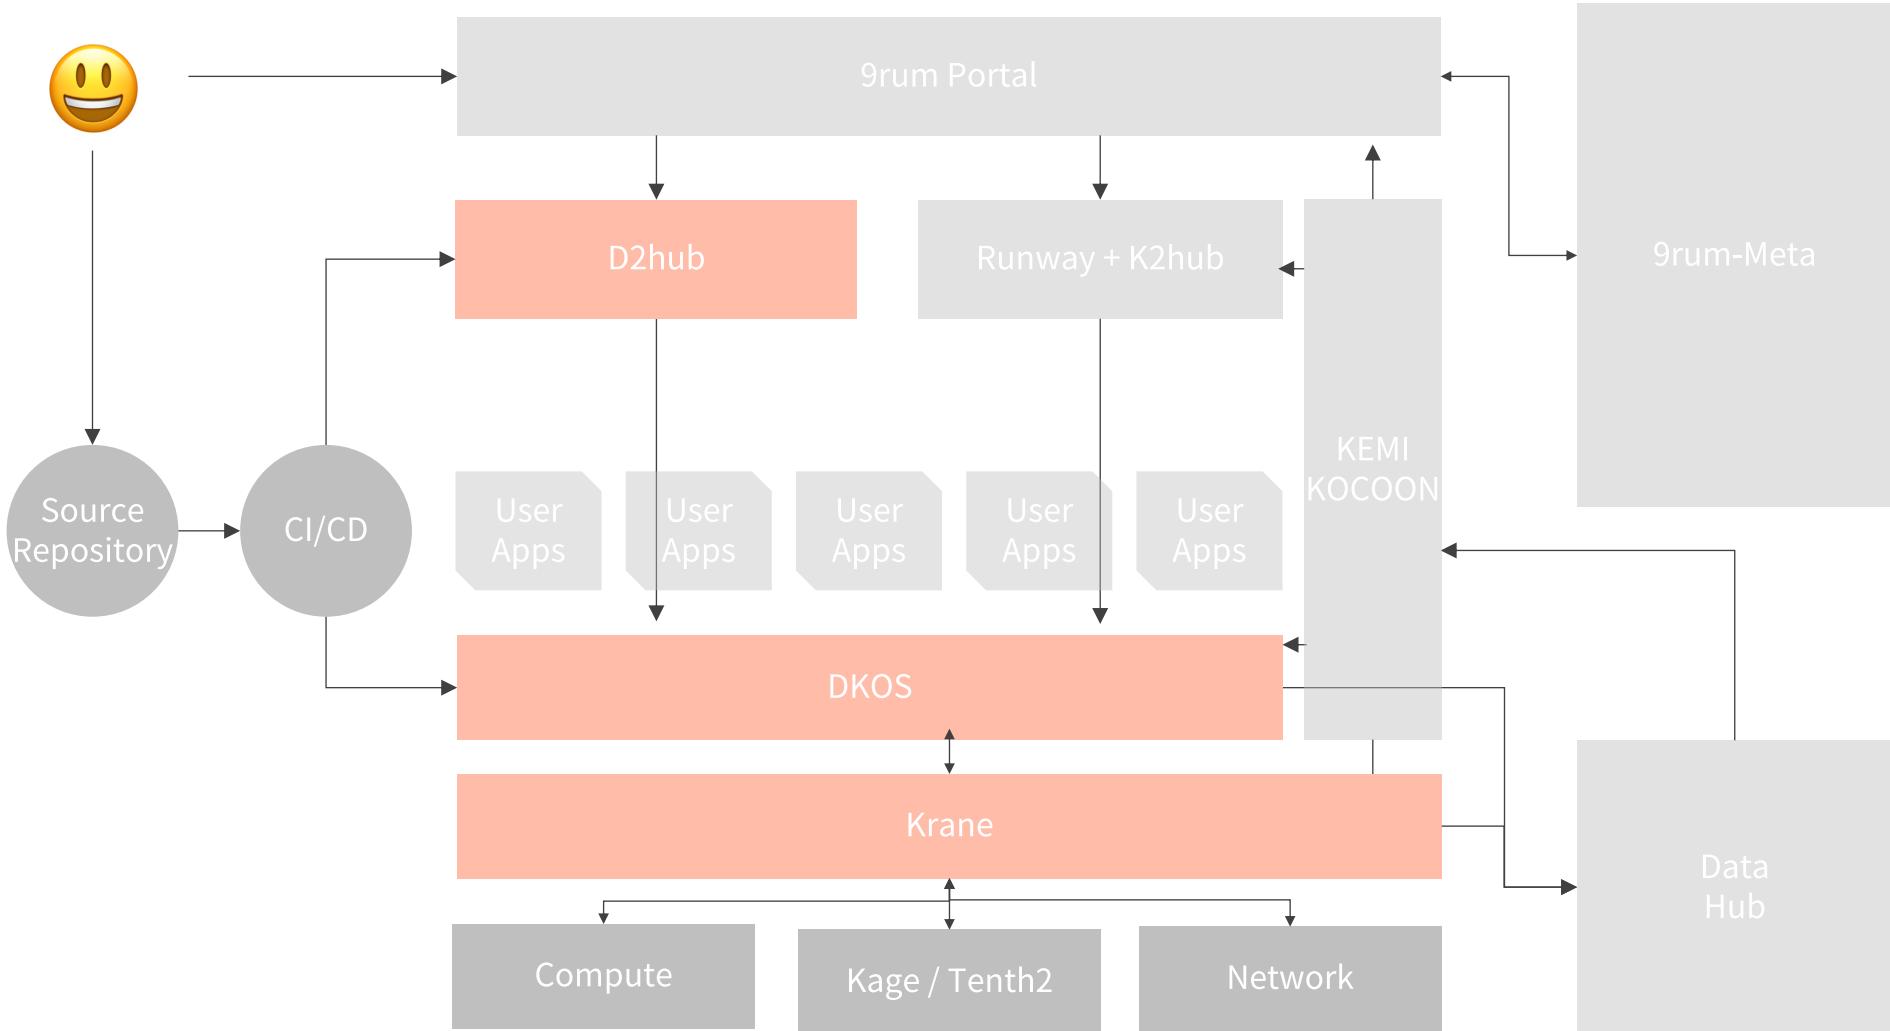
Kakao Cloud Architecture



Kakao Cloud Architecture



Our Product



Our Products



Infrastructure as a Service
OpenStack(stein)
Multiple Regions
L3 Network Model

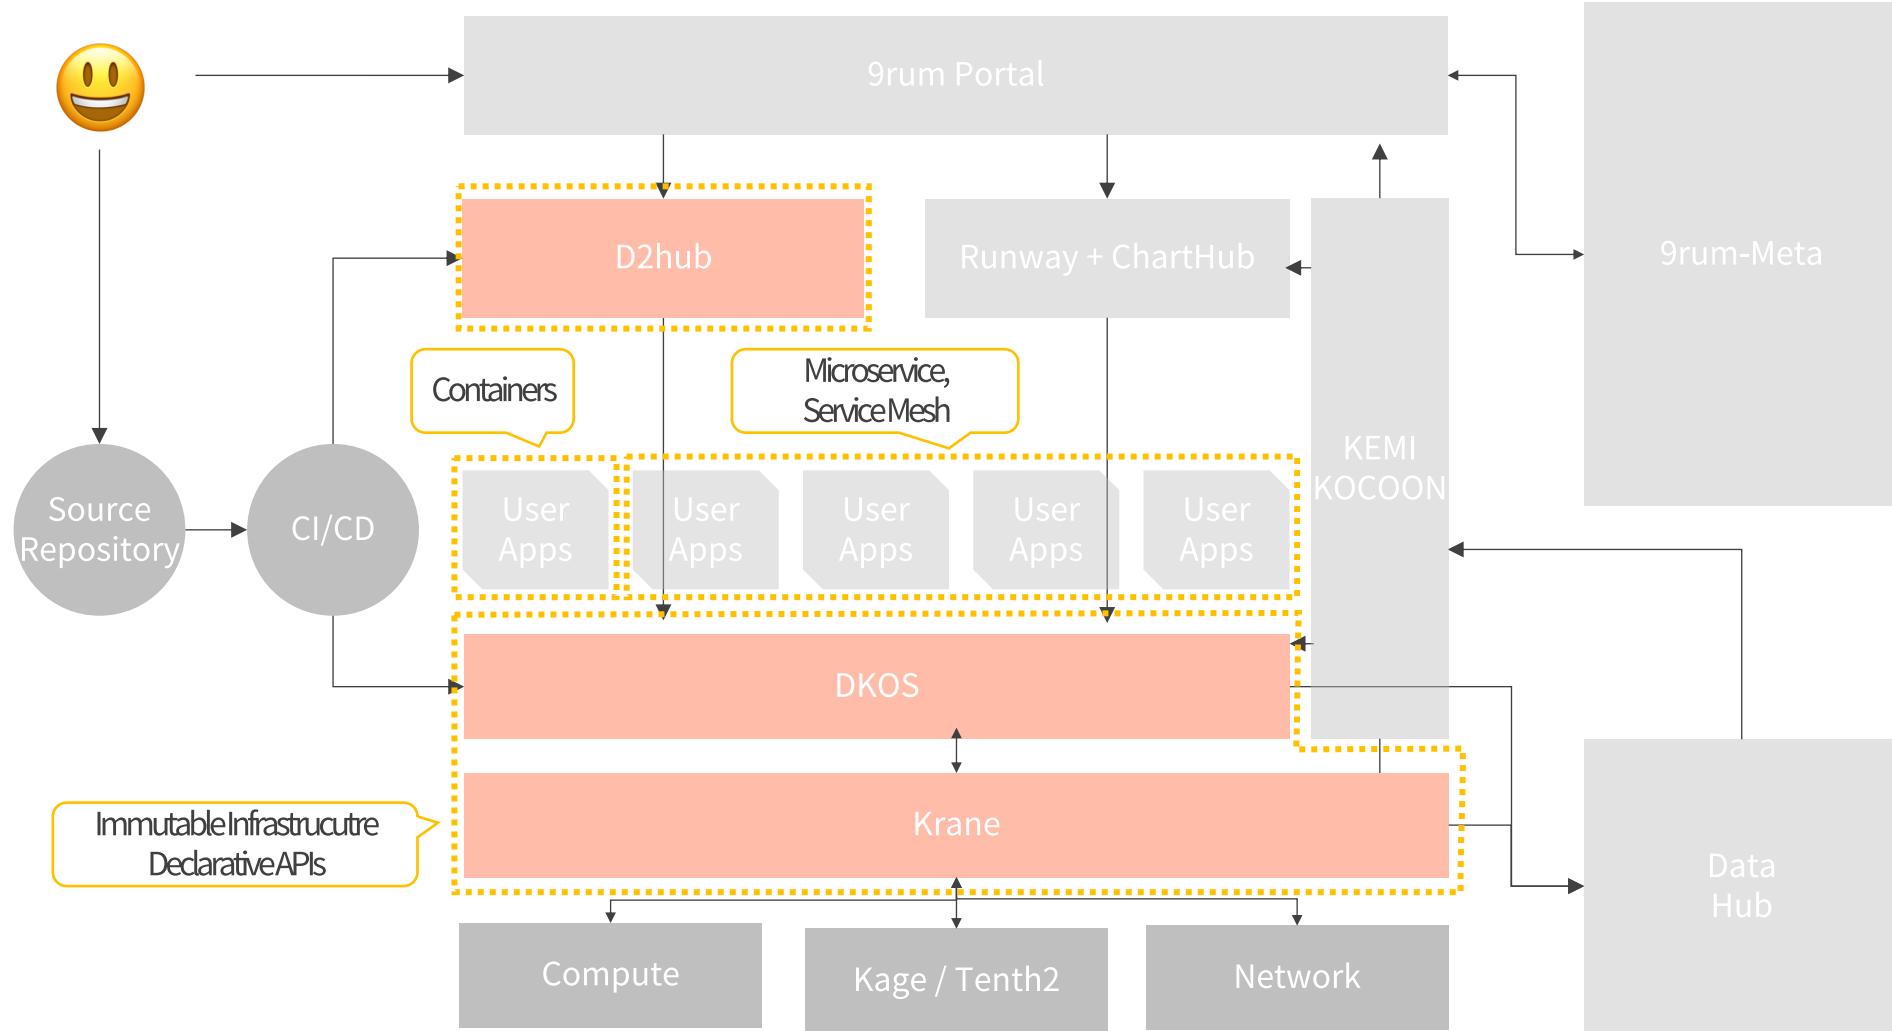


Container Cluster As A Service
Kubernetes
Custom Controller
Kakao Standard Service Platform

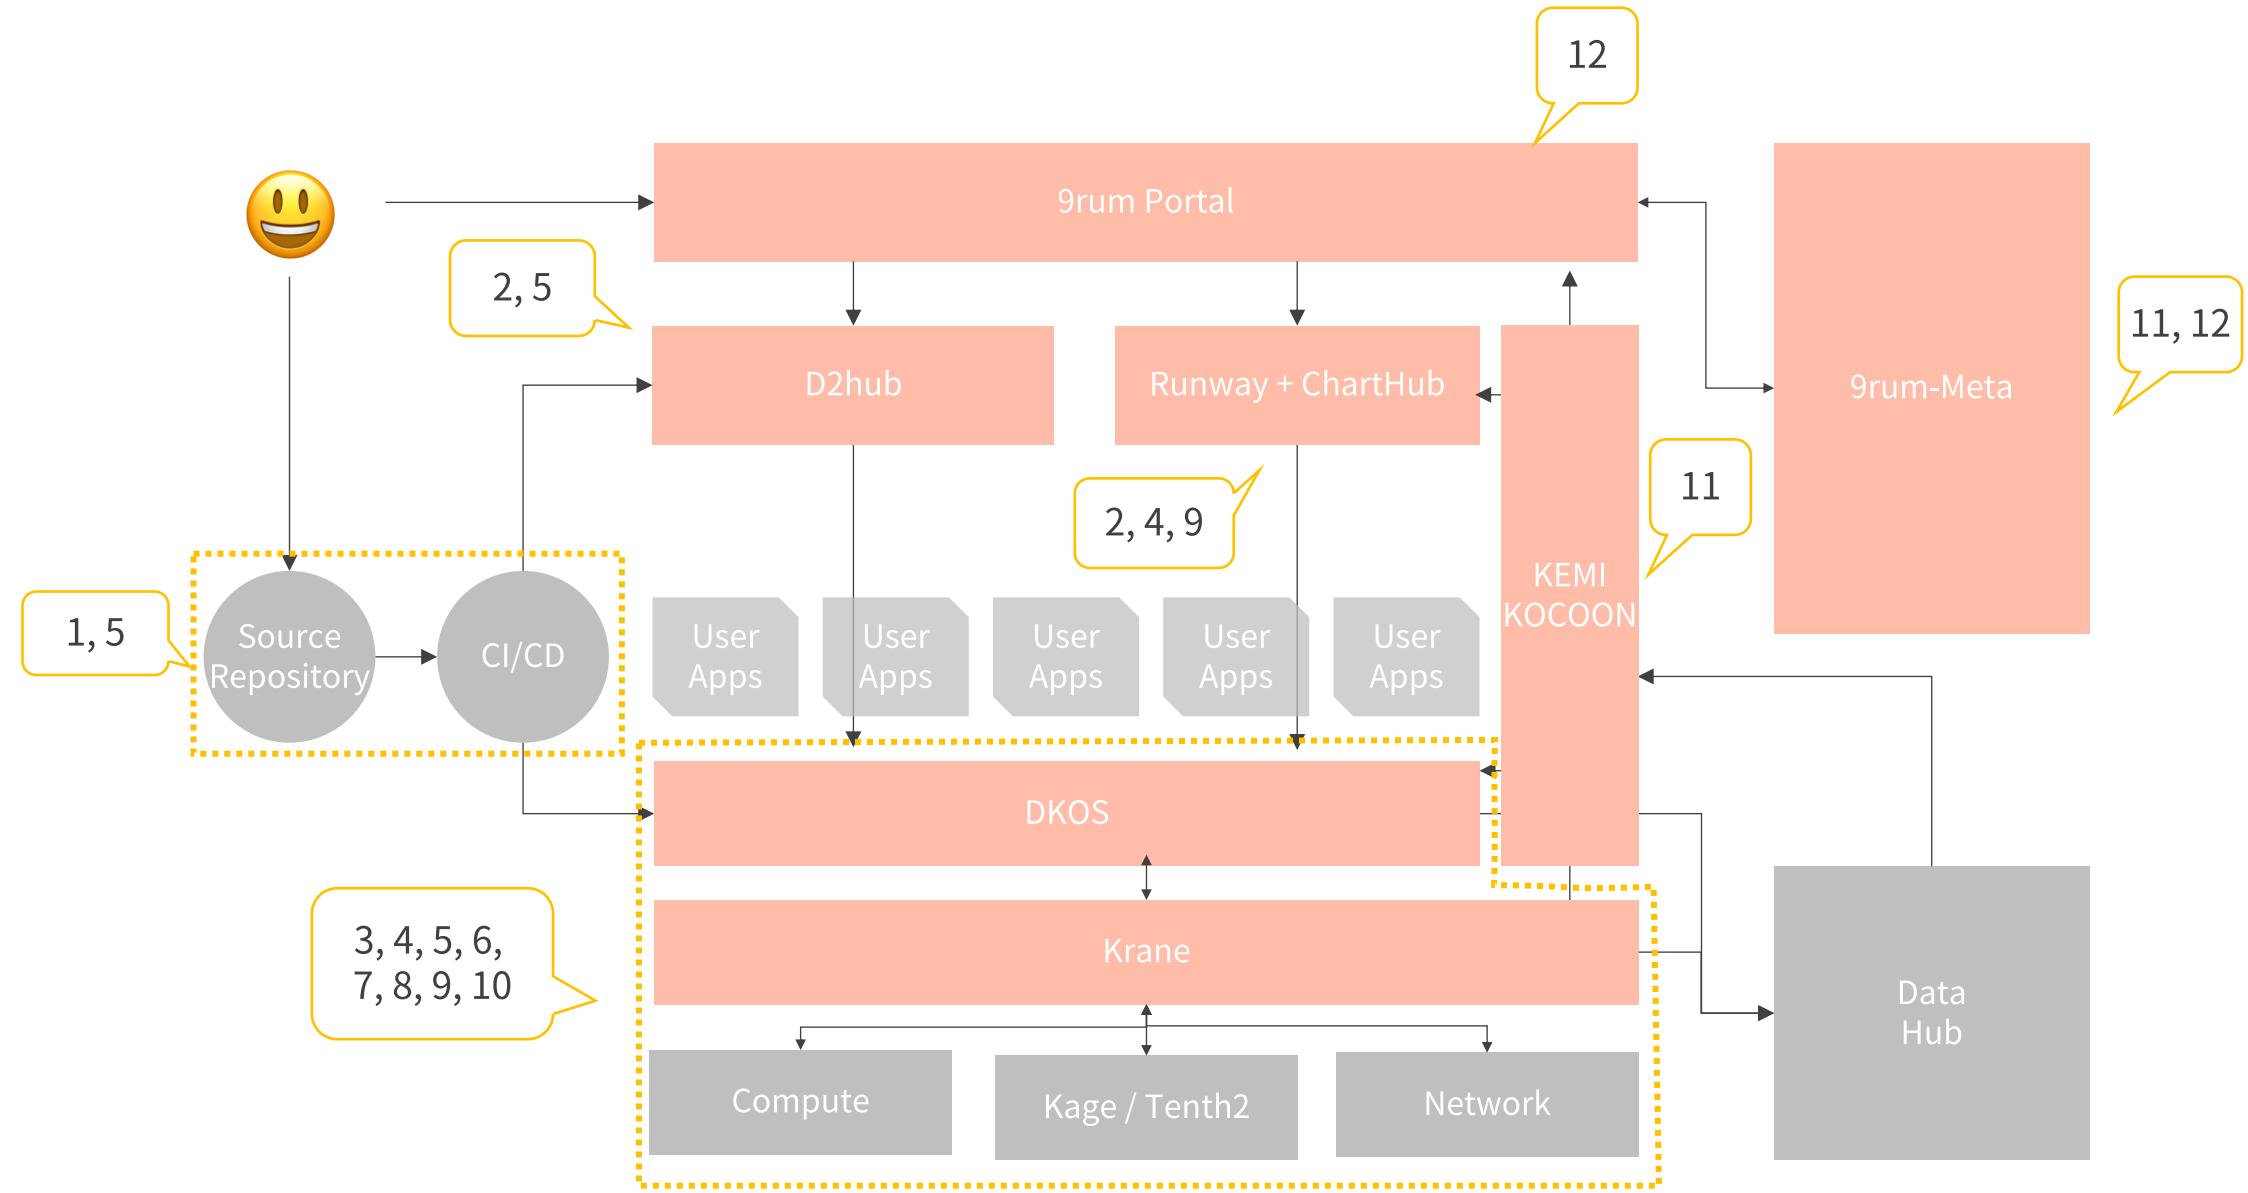


Container Image Registry
Docker Registry
Build and Inspect
Authentication and authorization

Support Cloud Native Technologies on Kakao Cloud



Support 12-factors on Kakao Cloud



03 How Kakao achieved to Cloud Native

Build Firm Foundation



Resource별 QoS 정책 적용

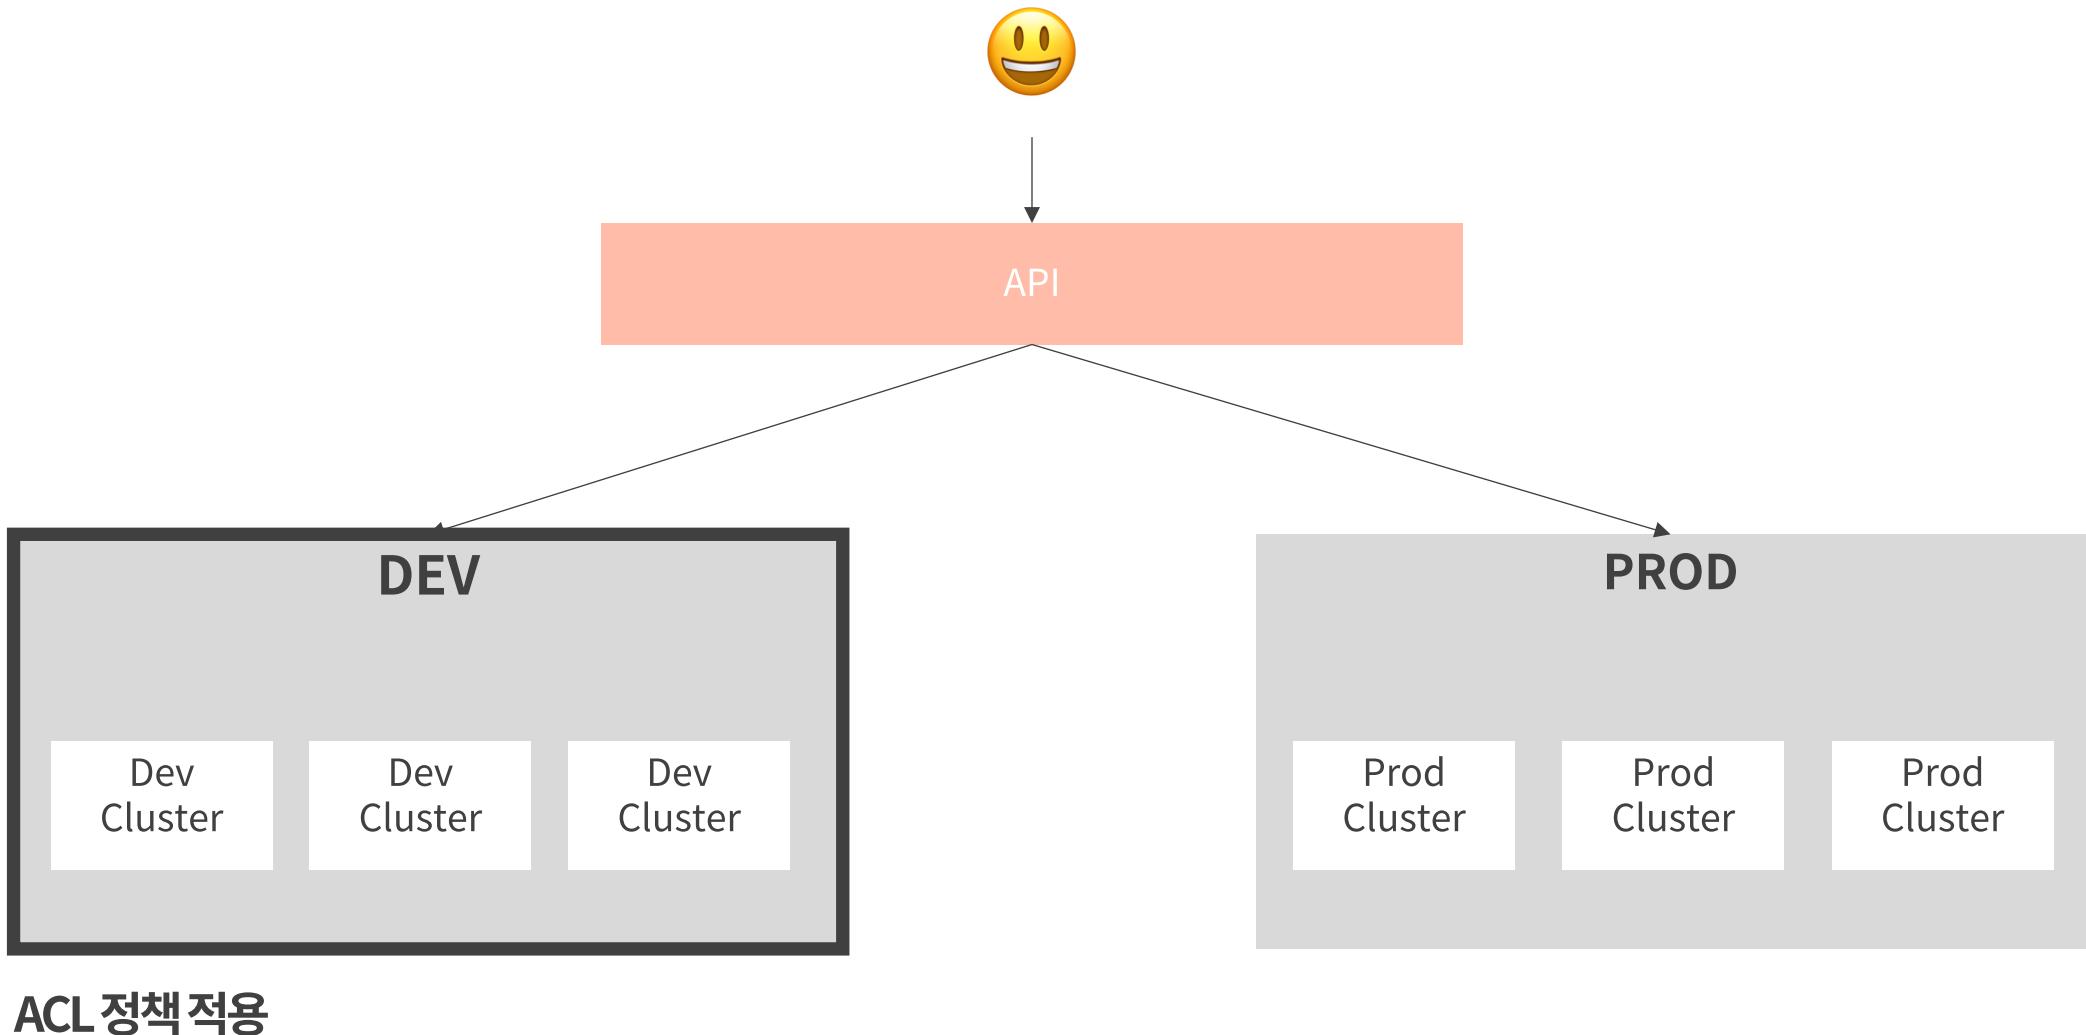


SLA에 최소 보장 Performance 명시



Performance Baseline을 올리기 위한 연구 개발

Kubernetes Cluster As A Service



Kubernetes Multitanency

One Big Cluster

Control Plane을 공유

자원의 효율적 사용 가능

User/Org 별 namespace로 구분

모든 사용자가 인프라 자원 공유

다른 사용자의 Pod와 보안 격리가 어려움



Multi Cluster

사용자 개별 클러스터 생성

클러스터 단위 자원 추가

사용자 인증 단위로 클러스터 구분

클러스터 단위 인프라 자원이 구분됨

클러스터 단위 보안 정책 적용 가능

Obstacle: Multi-Cluster Management

Web

사내 계정 시스템과 연동
참여 중인 클러스터만 표시
Kubernetes Token 기반으로 클러스터 Dashboard 접근

CLI

사내 자체 구축 Tool - **dkosctl**
인증 시스템 연동
참여 중인 프로젝트를 CLI 환경에서 context 전환

DKOSv3 Multi-Cluster View

The screenshot displays the DKOSv3 Multi-Cluster View interface. At the top, there is a blue header bar with the DKOS logo, a search bar, and user information for 'Lev.Jung (Member)'. Below the header, the interface is organized into sections for different clusters:

- krane**: This section contains two cluster cards. Both cards show a Kubernetes icon (a blue hexagon with a white steering wheel), the text "kubernetes version : v1.11.5 | v2.7.3", and a "Prod_Zone" button. Each card has four tabs at the bottom: "상세 보기" (selected), "KUBECTL 설정", "대쉬보드 이동", and "대쉬보드 토큰".
- lev**: This section contains one cluster card for "cloud-lev-dev". It shows a Kubernetes icon, the text "kubernetes version : v1.15.5 | v2.11.0", and a "Dev_Zone" button. The "lev test, dev cluster" text is also present. The card has the same four tabs at the bottom: "상세 보기" (selected), "KUBECTL 설정", "대쉬보드 이동", and "대쉬보드 토큰".
- prod**: This section contains one cluster card. It shows a Kubernetes icon, the text "kubernetes version : v1.15.5 | v2.11.0", and a "Prod_Zone" button. The card has the same four tabs at the bottom: "상세 보기" (selected), "KUBECTL 설정", "대쉬보드 이동", and "대쉬보드 토큰".

On the far right side of the interface, there is a vertical scroll bar and a blue circular button with a white plus sign (+).

DKOSv3 Cluster Detail View

DKOS
to cloud friendly way

Lev.Jung (Member)

cloud-lev-dev 클러스터 Dev_Zone

6 worker 1 ingress 3 master

방화벽 ACL Monitoring

Worker Nodes 로그 보기 All Nodes Healthy + Add Worker

data.cpu.user

agg.memory.used

-worker-1

ACTIVATED 225 / 3,900 m 111 / 3,600 MiB

-worker-2

ACTIVATED 225 / 3,900 m 111 / 3,600 MiB

-worker-3

ACTIVATED 325 / 3,900 m 131 / 3,600 MiB

-worker-4

ACTIVATED 275 / 3,900 m 172 / 3,600 MiB

dkosctl Multi-Cluster management

```
~> (* |krane-metric-context:default)dkosctl
DKOS controller
start with, "dkosctl configure"

To see help text, you can run:
  dkosctl [command] --help

Environment:
- $DKOSCTL_LDAP_ID # bypass login prompt
- $DKOSCTL_LDAP_PW
- $DKOSCTL_DKOSV3_HOST #  devel

Usage:
  dkosctl [command]

Available Commands:
  configure  configure
  console    openstack console
  create     create
  dashboard  open kubernetes dashboard
  delete    delete
  get        get
  help       Help about any command
  ns         namespace
  openstack  generate openstack env(osenv)
  project   openstack project
  quota     openstack quota
  server    openstack server
  ssh       ssh
  version   version
  volume    openstack volume

Flags:
  --debug   debug mode
  -h, --help  help for dkosctl

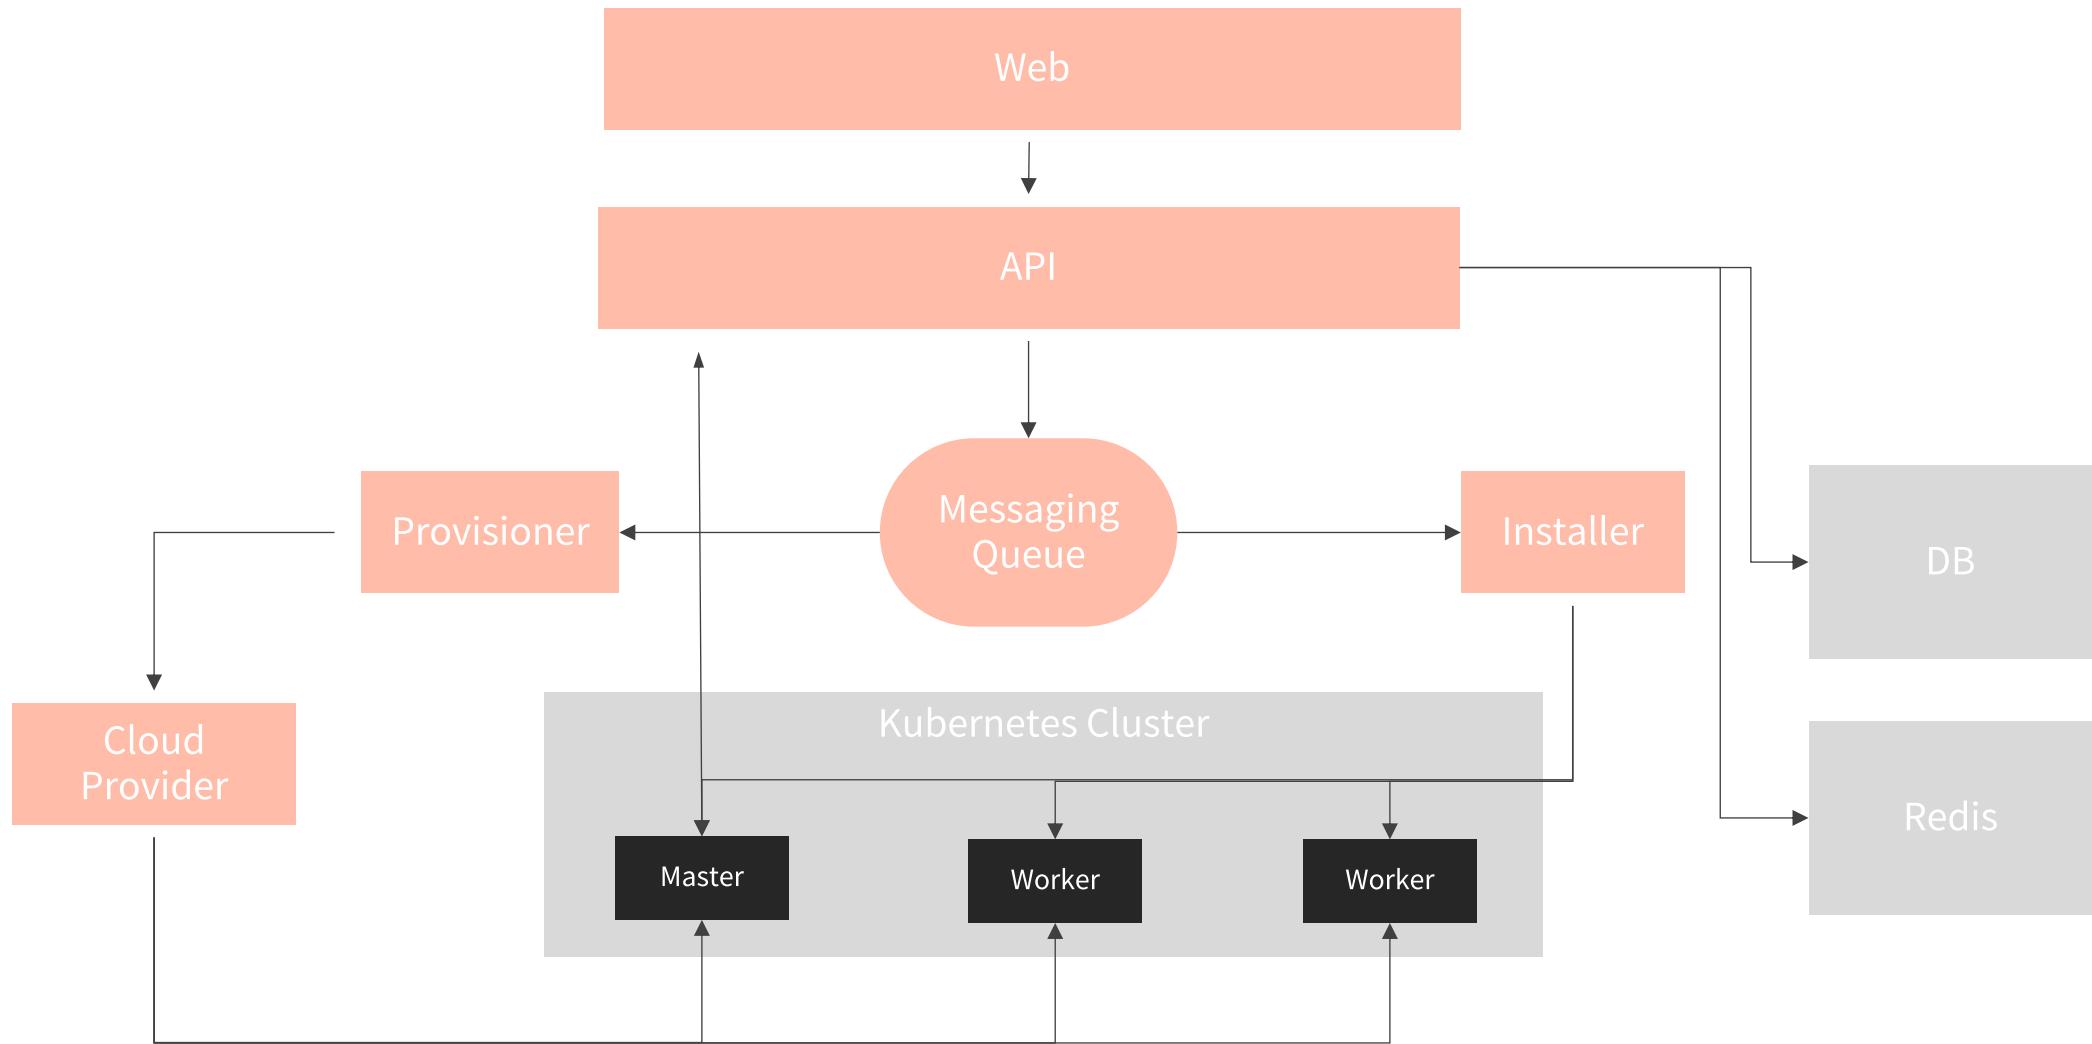
Use "dkosctl [command] --help" for more information about a command.
```

dkoctl cli 제공 가능

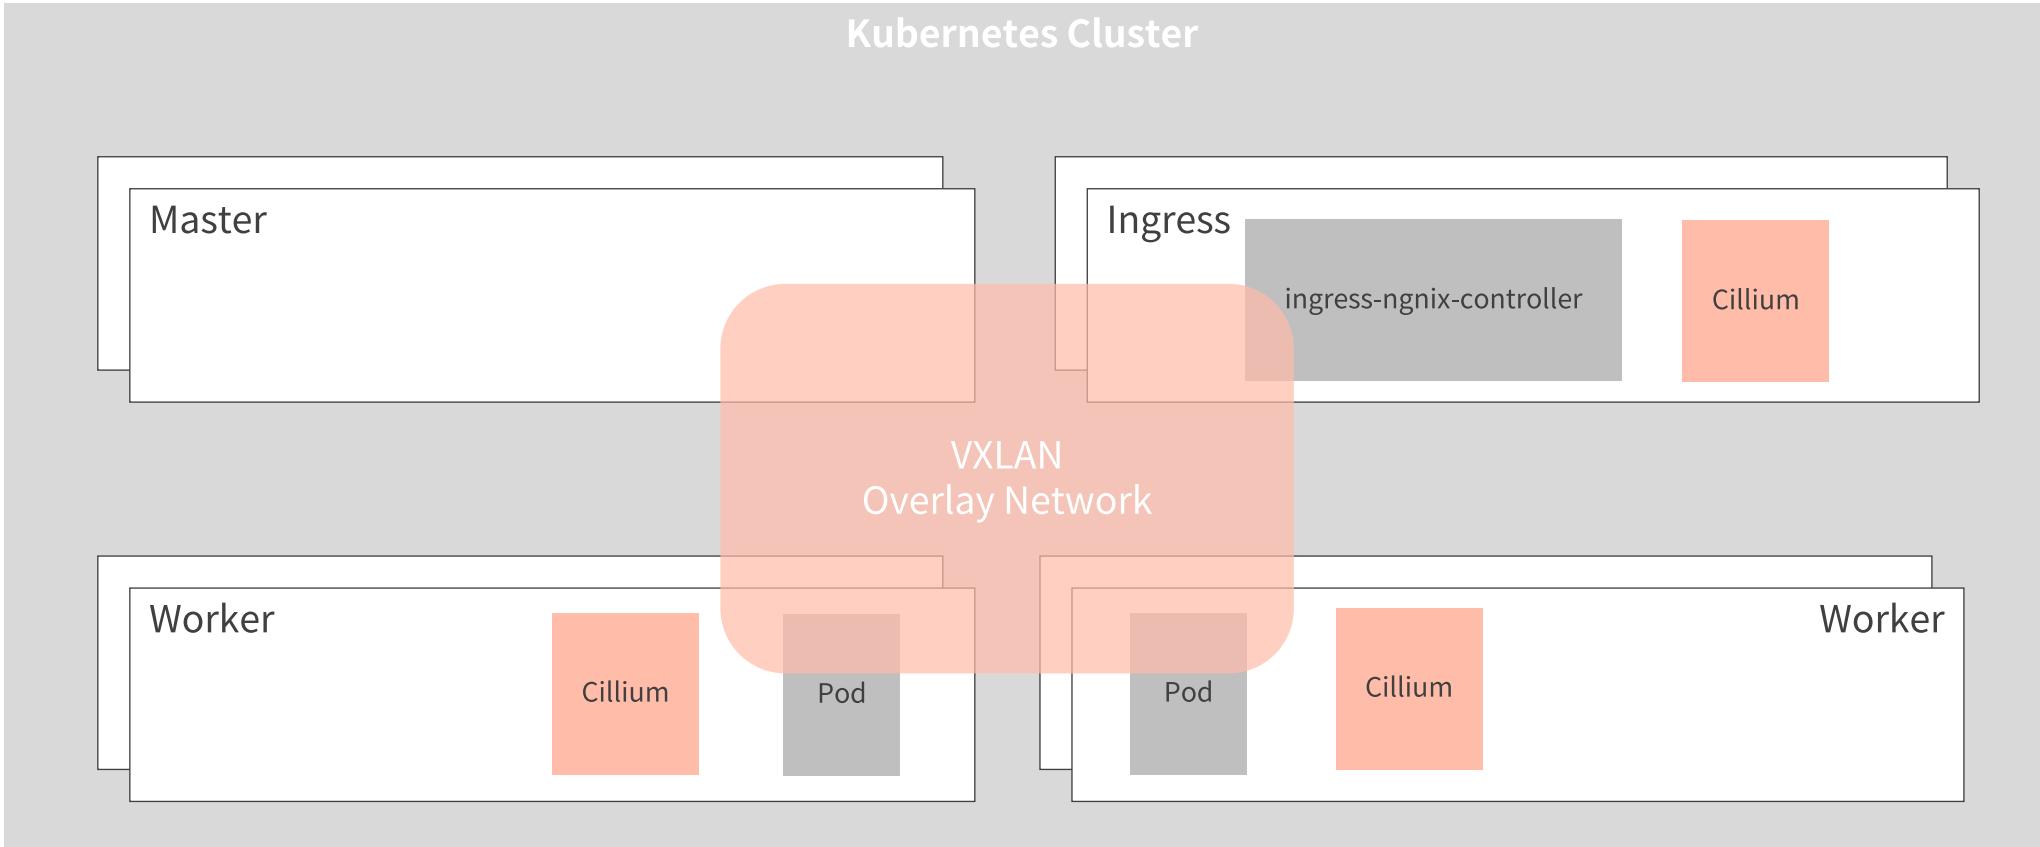
ID	NAME	KUBERNETES	KUBESPRAY	IMS(ROOT)	CREATED	OPENSTACK	DESCRIPTION
4910	cloud-lev-dev	v1.15.11 v1.15.5 v1.15.5 v1.11.5 v1.11.5 v1.11.5	v2.11.1 v2.11.0 v2.11.0 v2.7.3 v2.7.3 v2.7.3		true false false true true true		lev test, dev c
>							
> □	6/6						

dkosctl multicloud 선택 예시

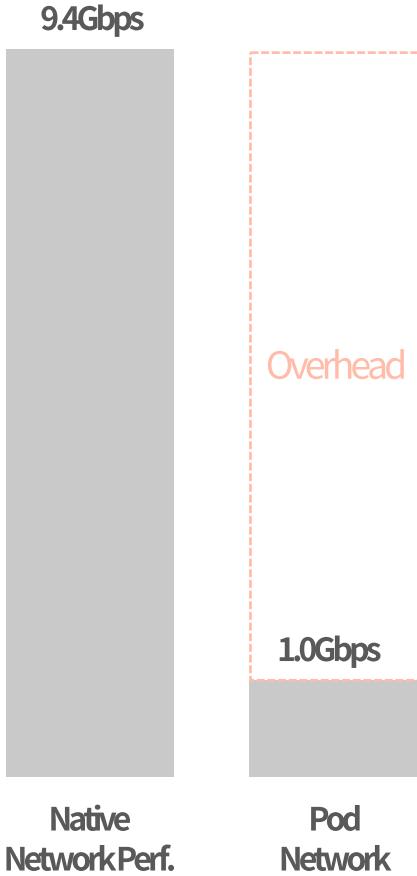
DKOSv3 Architecture Overview



DKOSv3 Cluster



Obstacle: Vxlan Performance



Tunnel Overhead

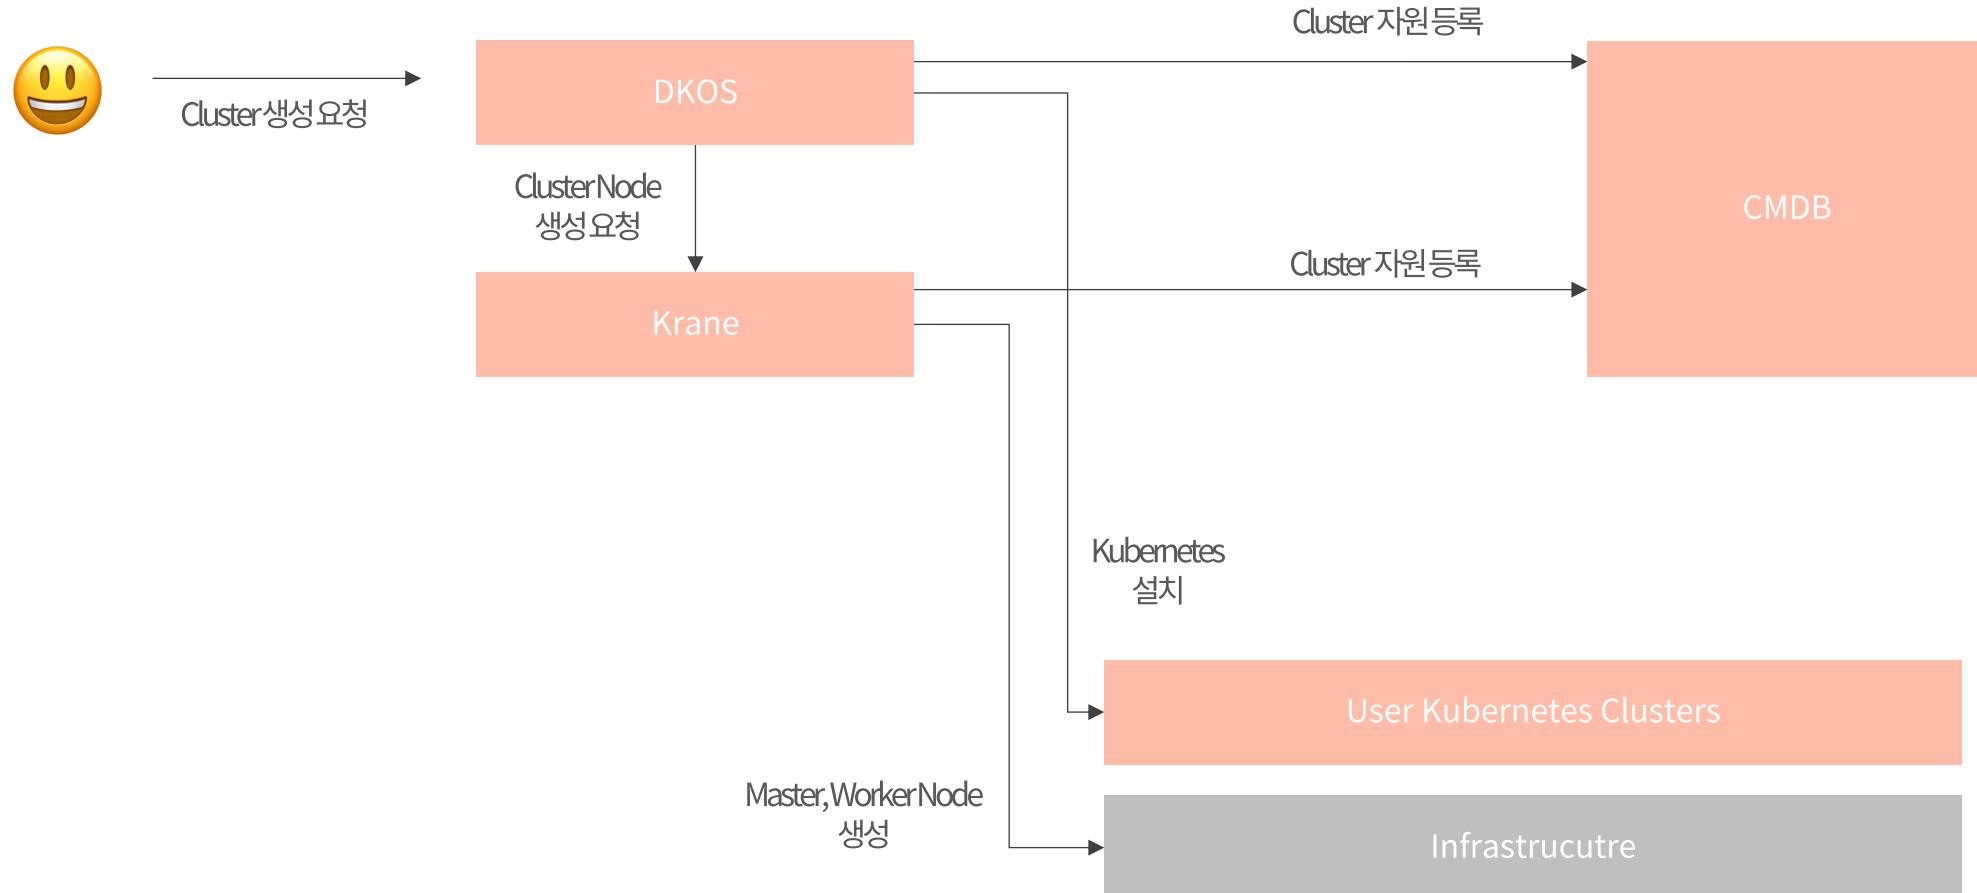
Cilium Datapath 구조 분석

VxLAN encapsulation / decapsulation에 의한
오버헤드 발생인지

Pod간 Overlay Network 필요하지 않는 경우 HostNetworkMode

SR-IOV와 같이 NIC Offloading 기능 활용 개발 방향 수립

Code to App: Cluster



Obstacle: Infra Resources

Image Based Infra

Master, Worker, Ingress 모두 Krane을 기반으로 생성하는 VM

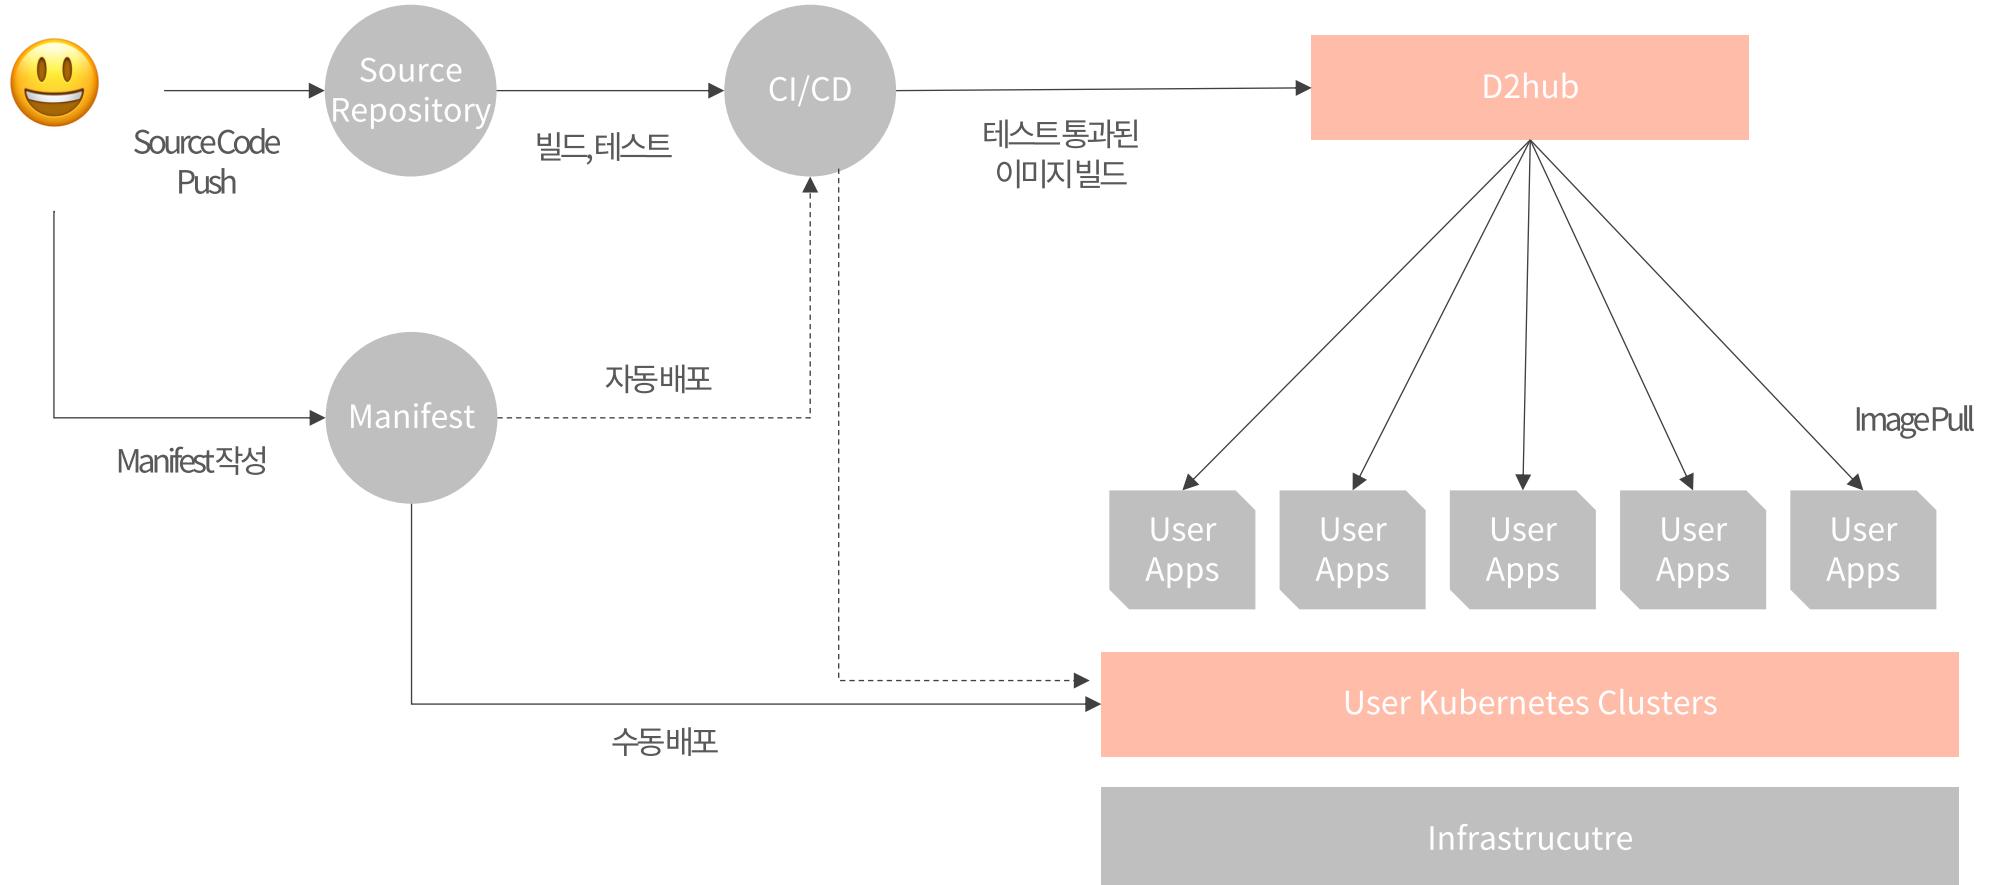
VM 이미지는 인프라팀에서 사내 보안 규정을 만족하는 이미지(Base Image) 생성

Base Image 기반으로 매주 자동으로 빌드

노드 내부 설정 오염 또는 소프트웨어 이슈 해결은 Rebuild or Re-Create

보안 패치, 커널 업그레이드가 필요한 경우도 Rebuild or Re-Create

Code to App: Deployment



Obstacle: Container Image

D2hub

사내 Github Enterprise와 연동하여 Repo의 Branch, Tag 기반으로 Docker Image 빌드 기능 지원

LDAP과 연동하여 인증과 권한 관리

사내 보안 도구와 연계하여 Image 내 알려진 취약성 검사 가능

인가된 클러스터에 이미지를 자동 배포 기능 제공

Obstacle: Manifest

K2hub

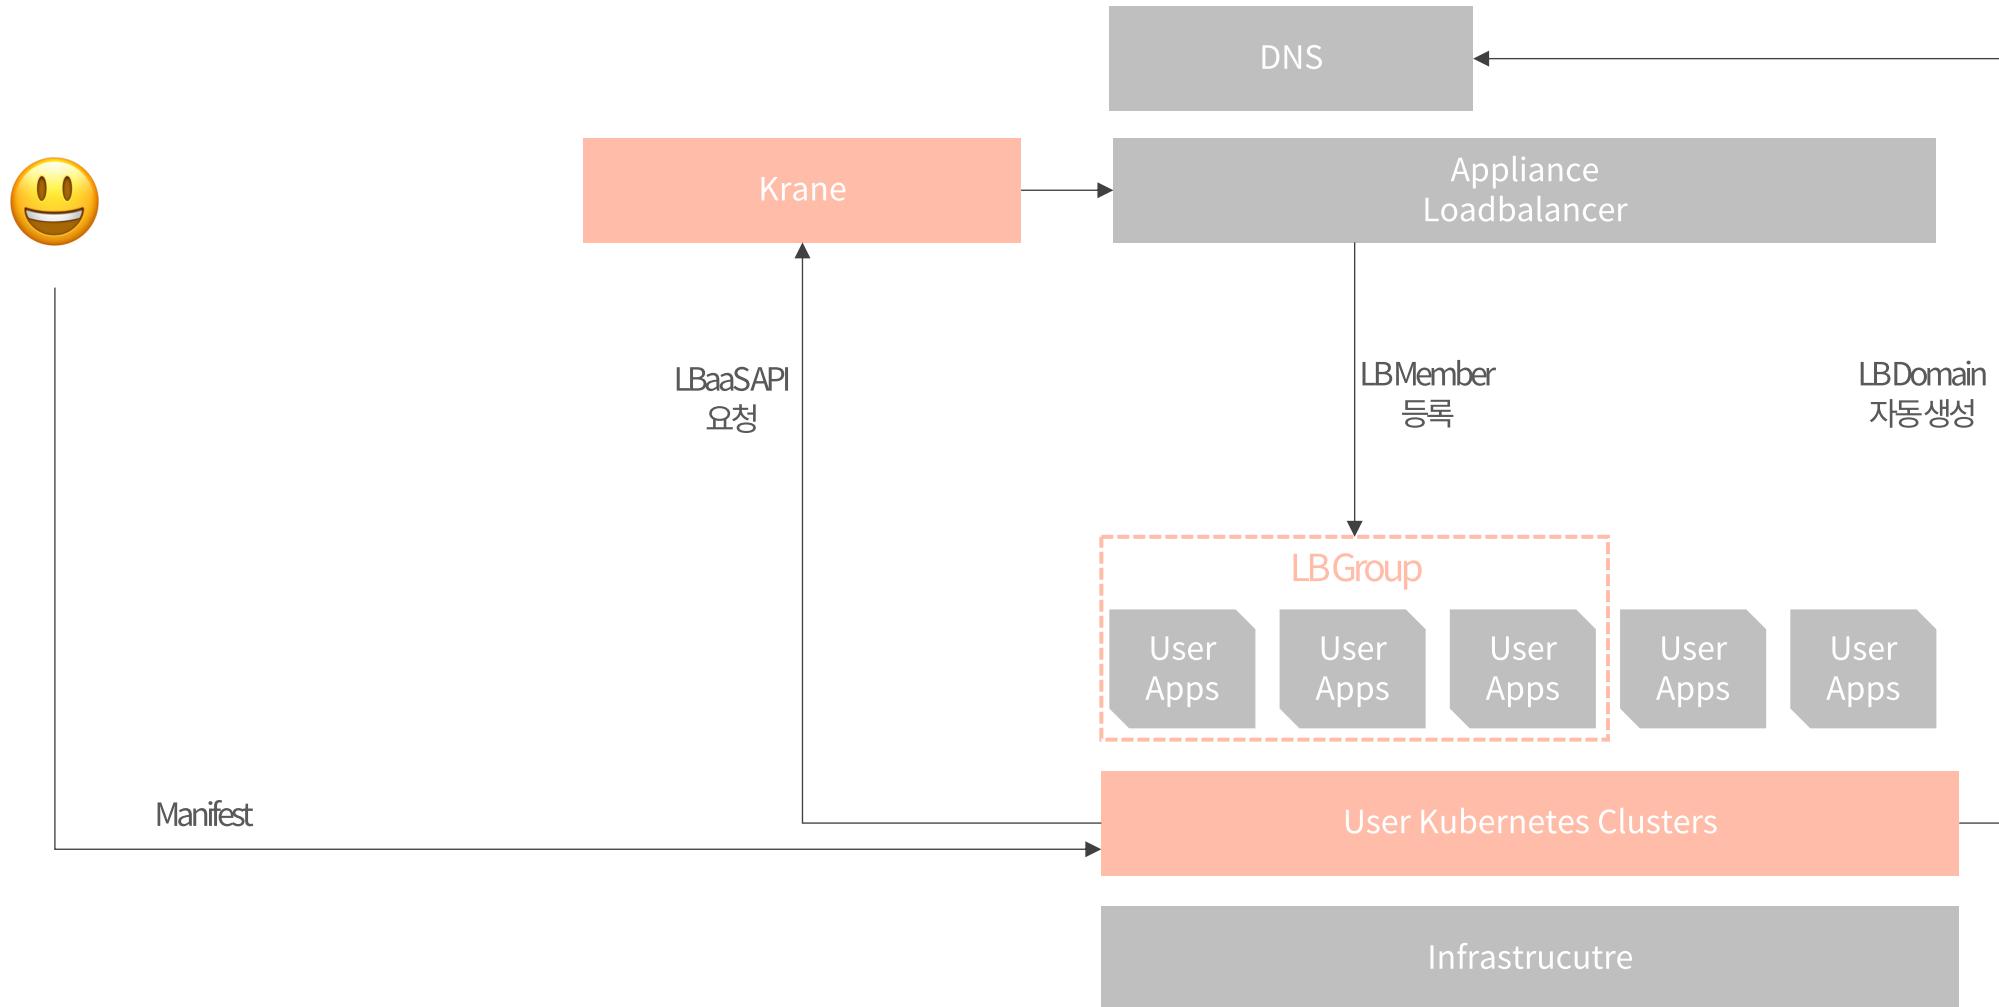
Kakao Helm Chart museum

사내에서 관리하는 Application / Platform의 Chart 사용 가능

잘 알려진 외부 Chart로 제공

Web 기반으로 인가된 클러스터에 Chart 배포 기능 제공

Code to App: Connect to LoadBalancer



Obstacle: DSR

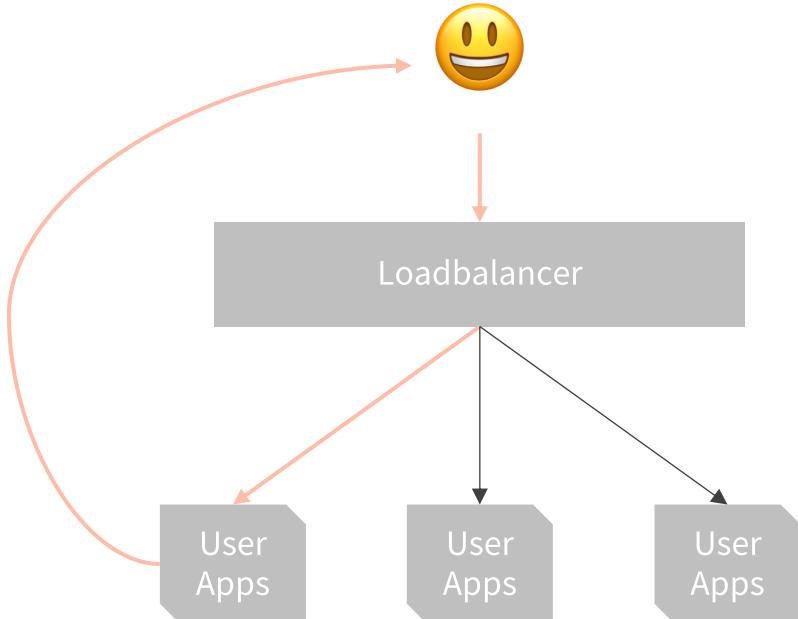
Direct Server Return(DSR)

카카오 사내 LB 기본 정책

LB와 Member 간의 IP-in-IP 터널이 필요

DSR-Manger 구현

자동으로 Ingress Node, Worker Node에서 DSR 설정 역할



Obstacle: LB Controller

Custom LB Controller

OpenStack이 관리하지 않는 PM의 경우도 LB Member로 참여 가능

Node의 Label 기반으로 LB Member Selection

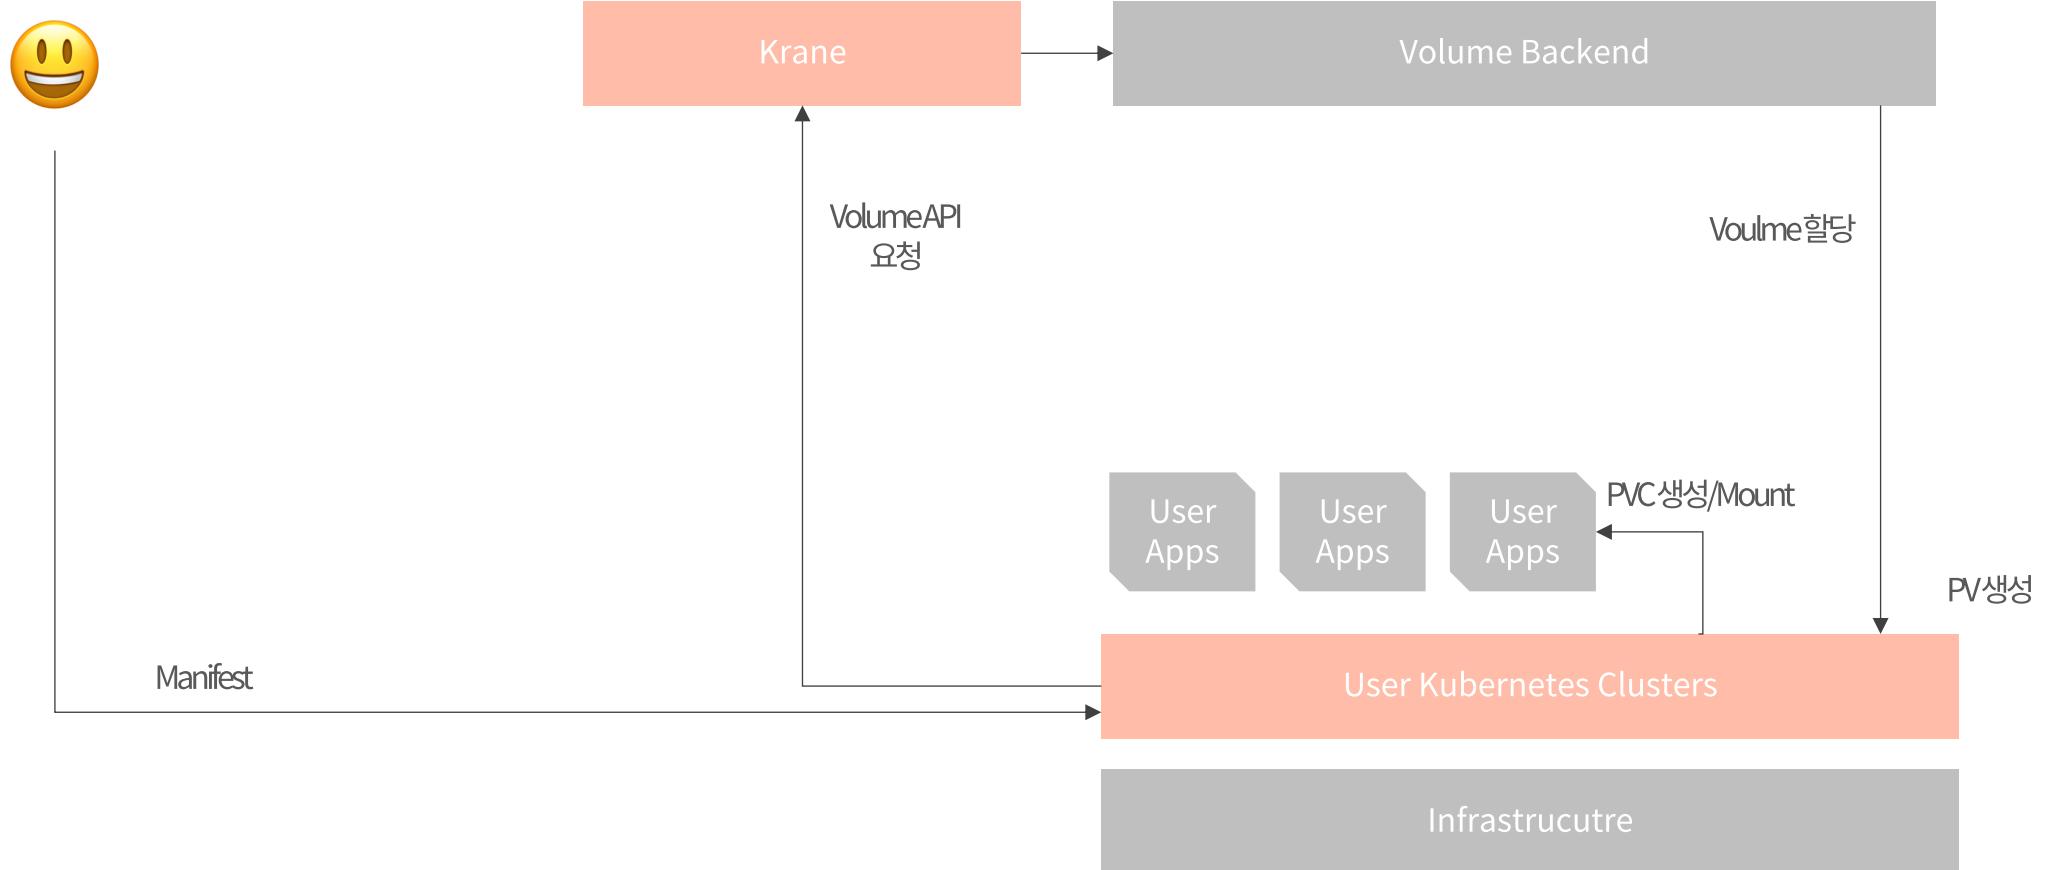
Dev/Prod Cluster 별로 다른 LB Provider 선택 가능

Obstacle: Domain

inhouse-dns-controller

Ingress 설정에 특정 형식으로 Host 설정을 한 경우 자동으로 DNS에 domain 등록

Code to App: Connect to Volume

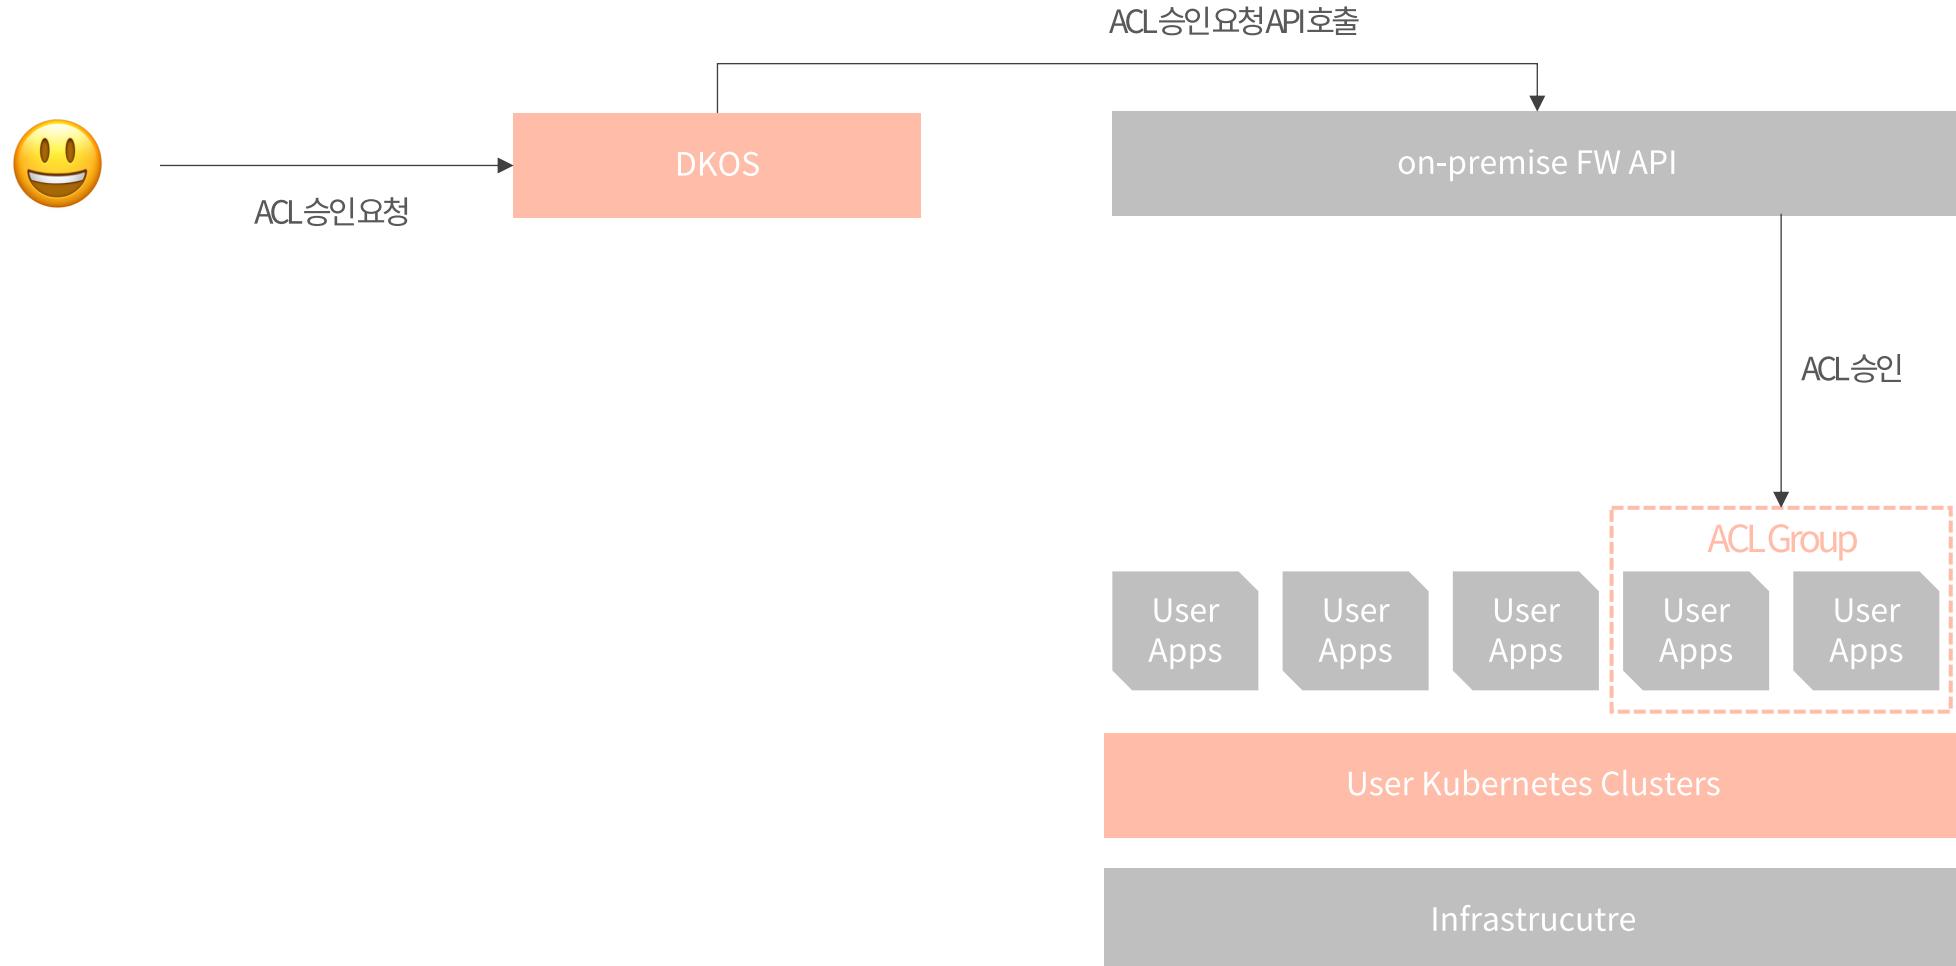


Obstacle: Cinder Controller

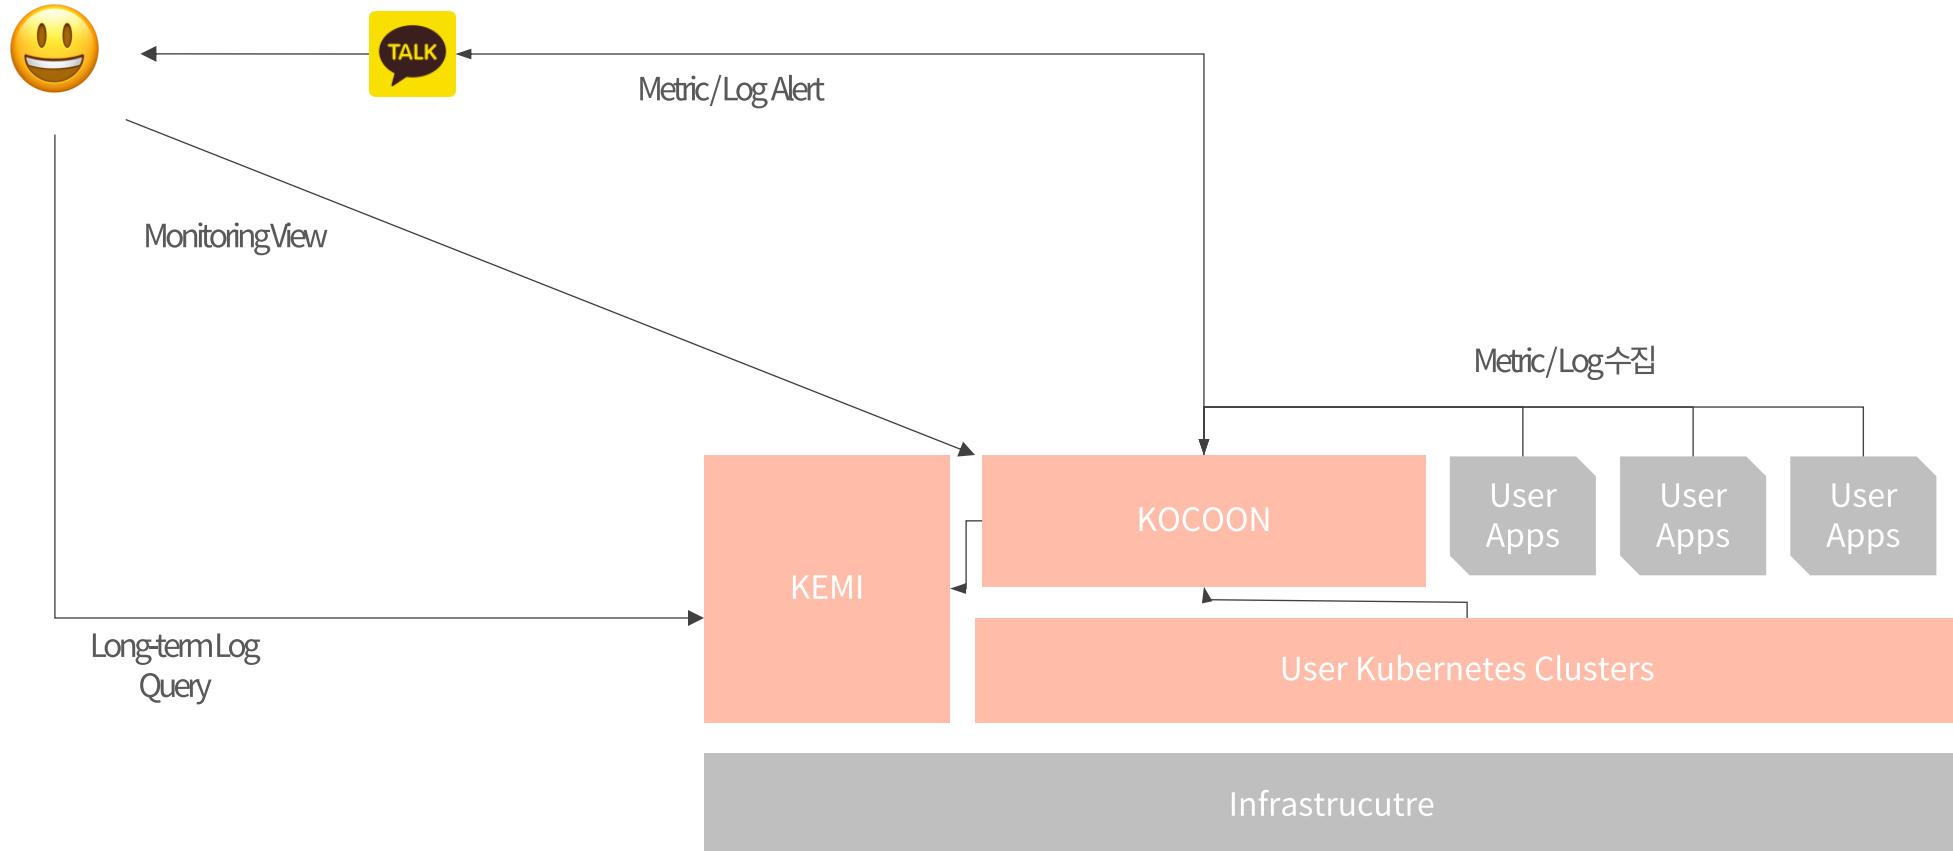
CSI Cinder

Cinder Controller를 사용하기 위해 Kubernetes Cloud Provider가 OpenStack API에 주는 부하가 상당
부하를 줄이기 위해 CSI(Container Storage Interface)를 활용한 새로운 Storage Controller를 구현

Code to App: Connect to ACL



Code to App: Monitoring



Obstacle: Resource Monitoring

KOCOON

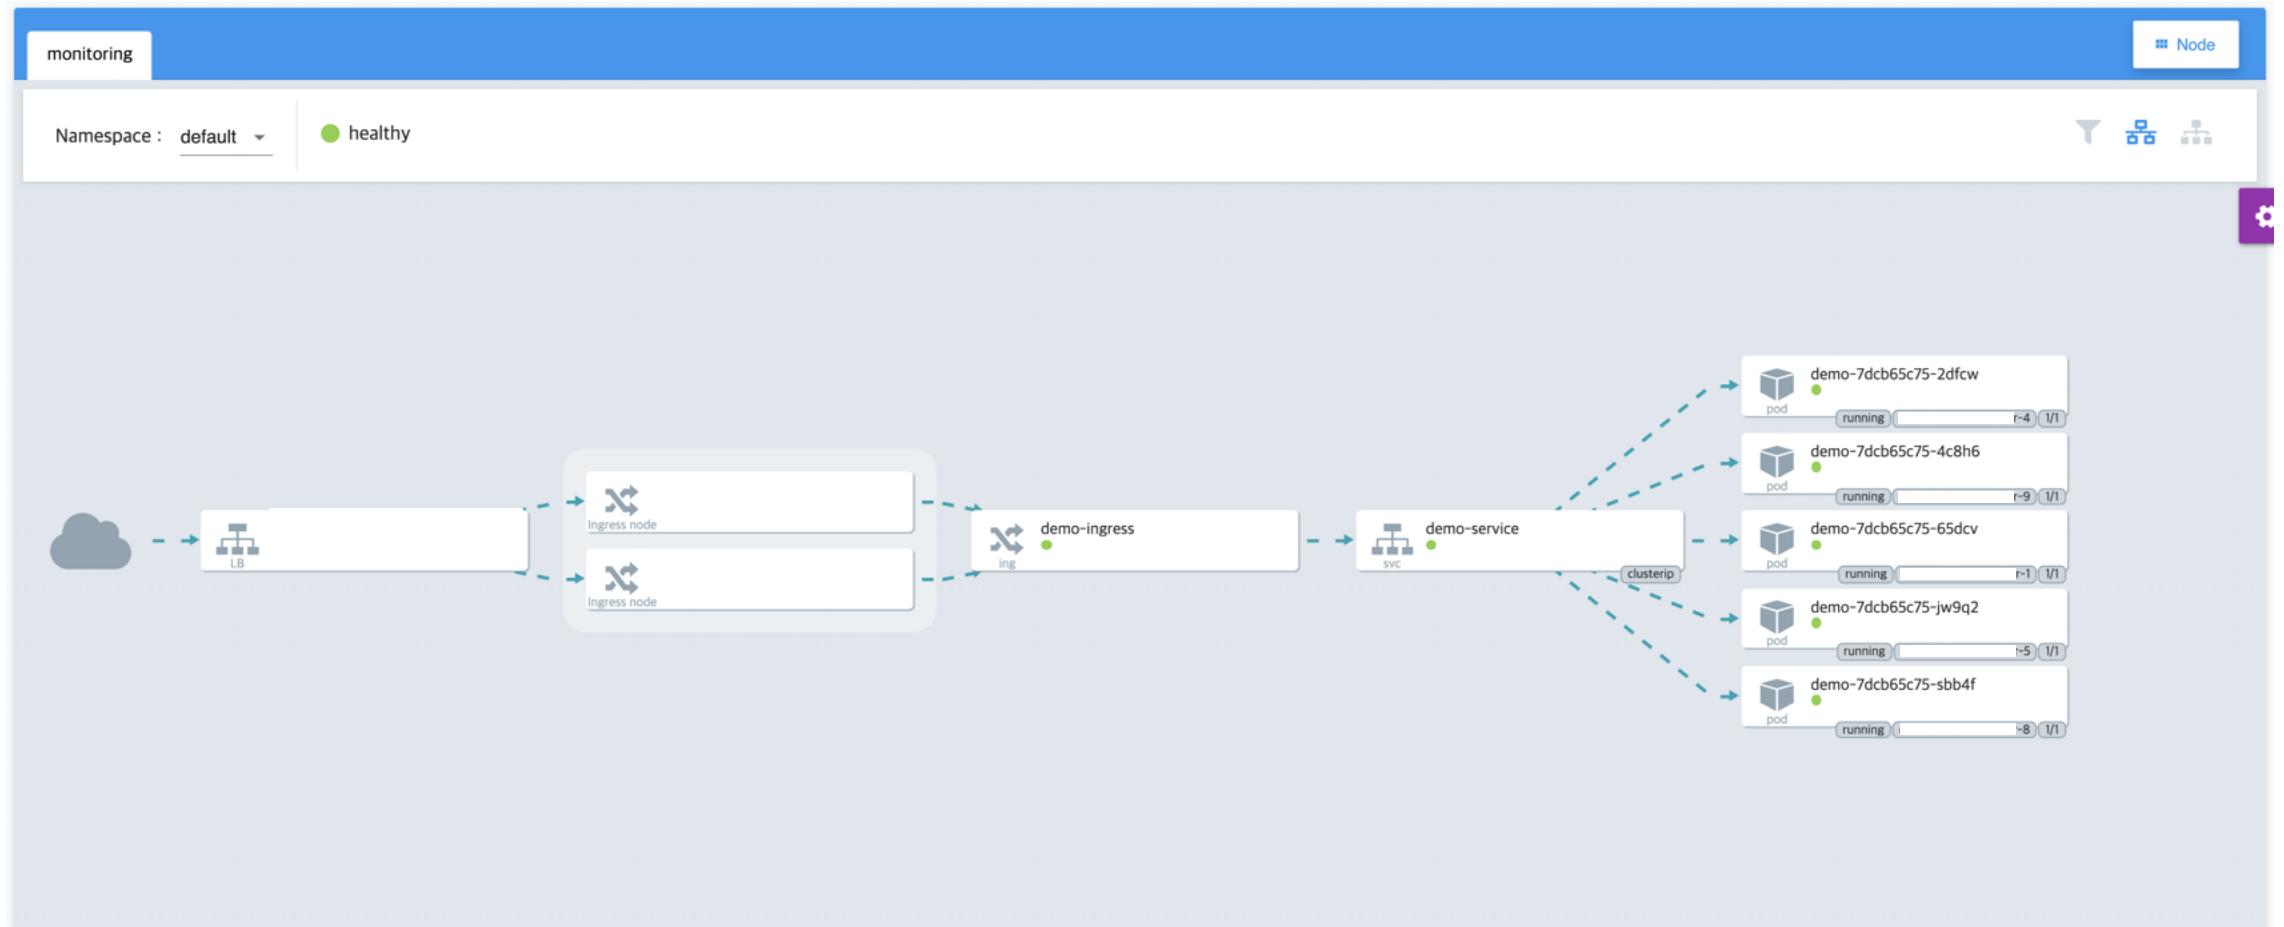
서비스 클러스터 별로 모니터링 리소스를 분리

Prometheus 기반의 메트릭 수집

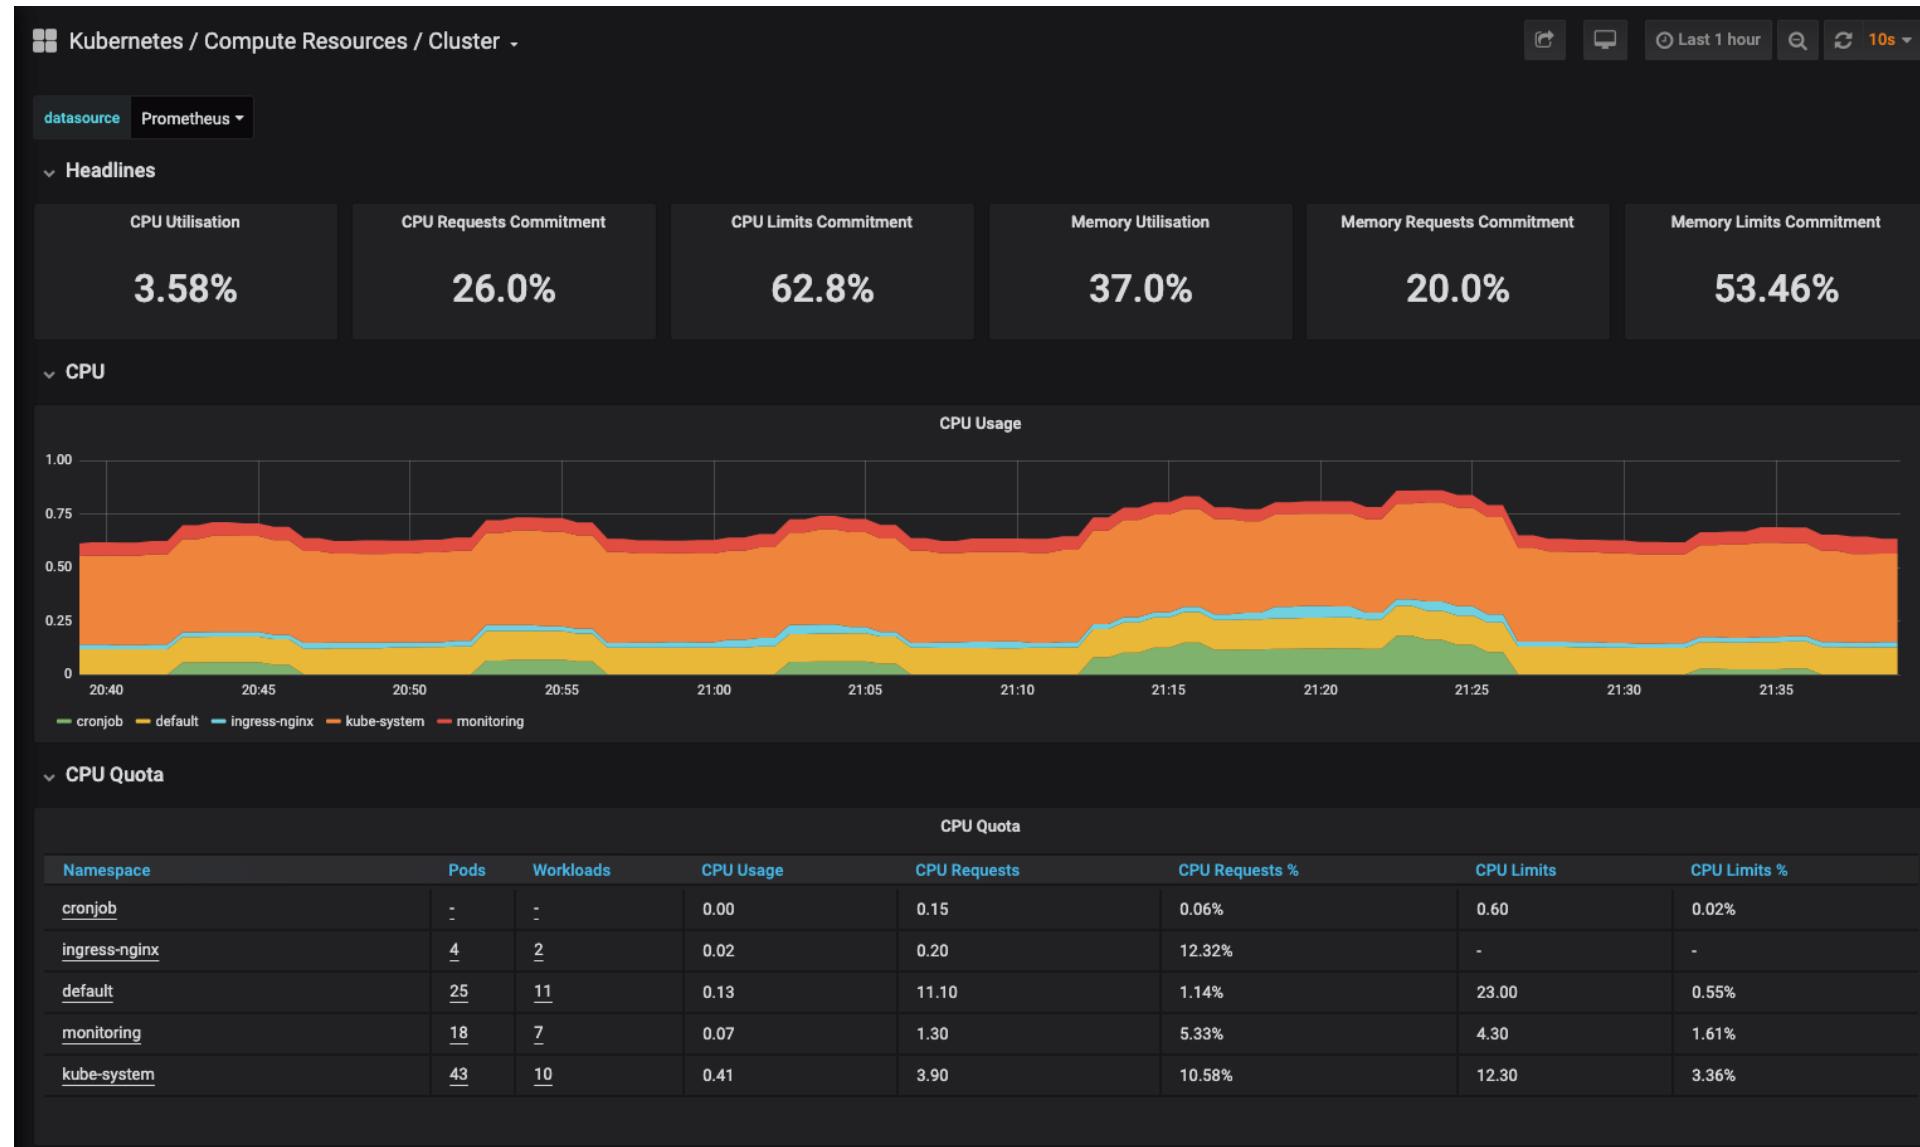
KOCOON-Hermes(Fluentd 기반)으로 Log Routing

KOCOON-CUPIDO로 Alert을 카톡 전송

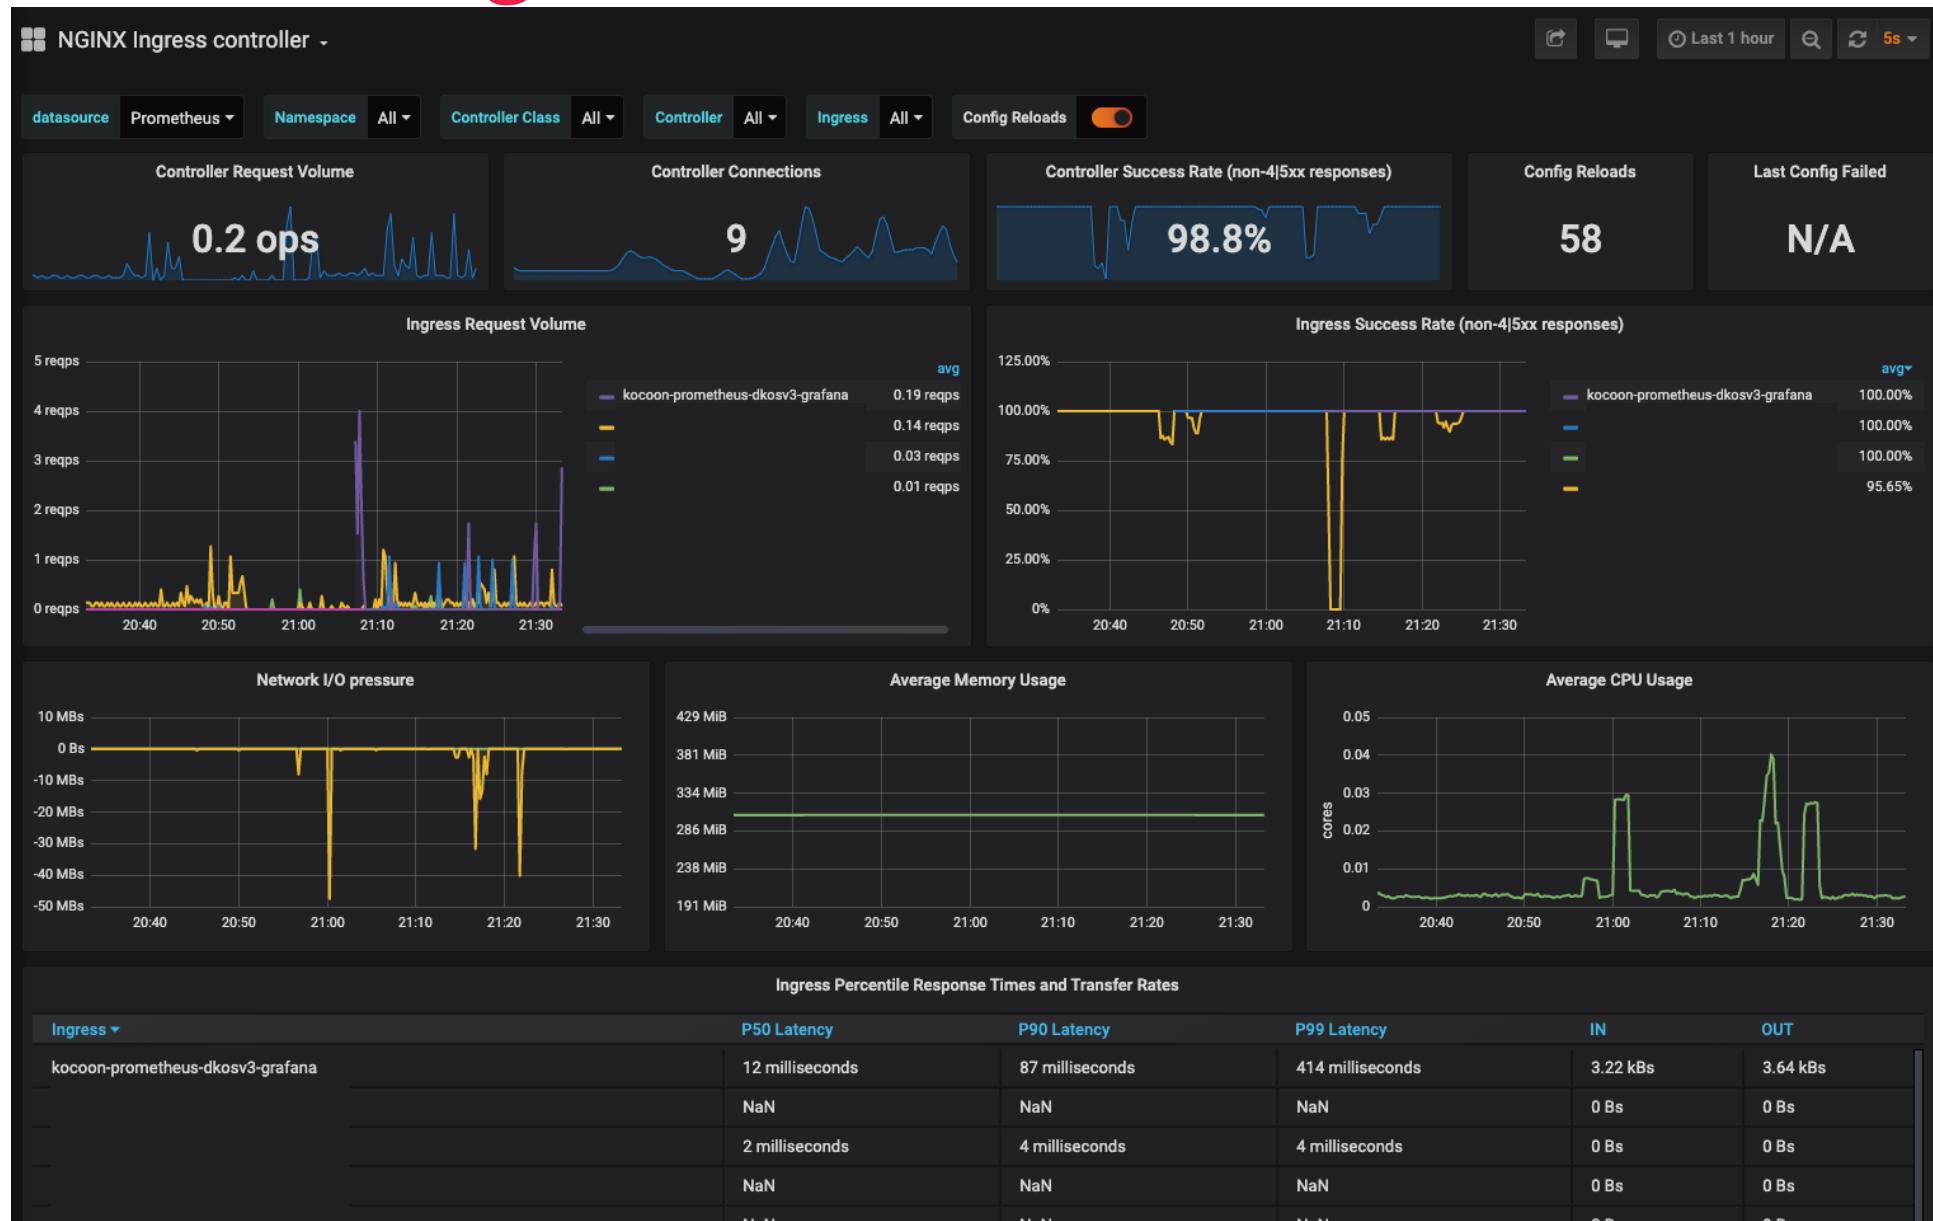
DKOSv3 Monitoring View



DKOSv3 Monitoring View



DKOSv3 Monitoring View

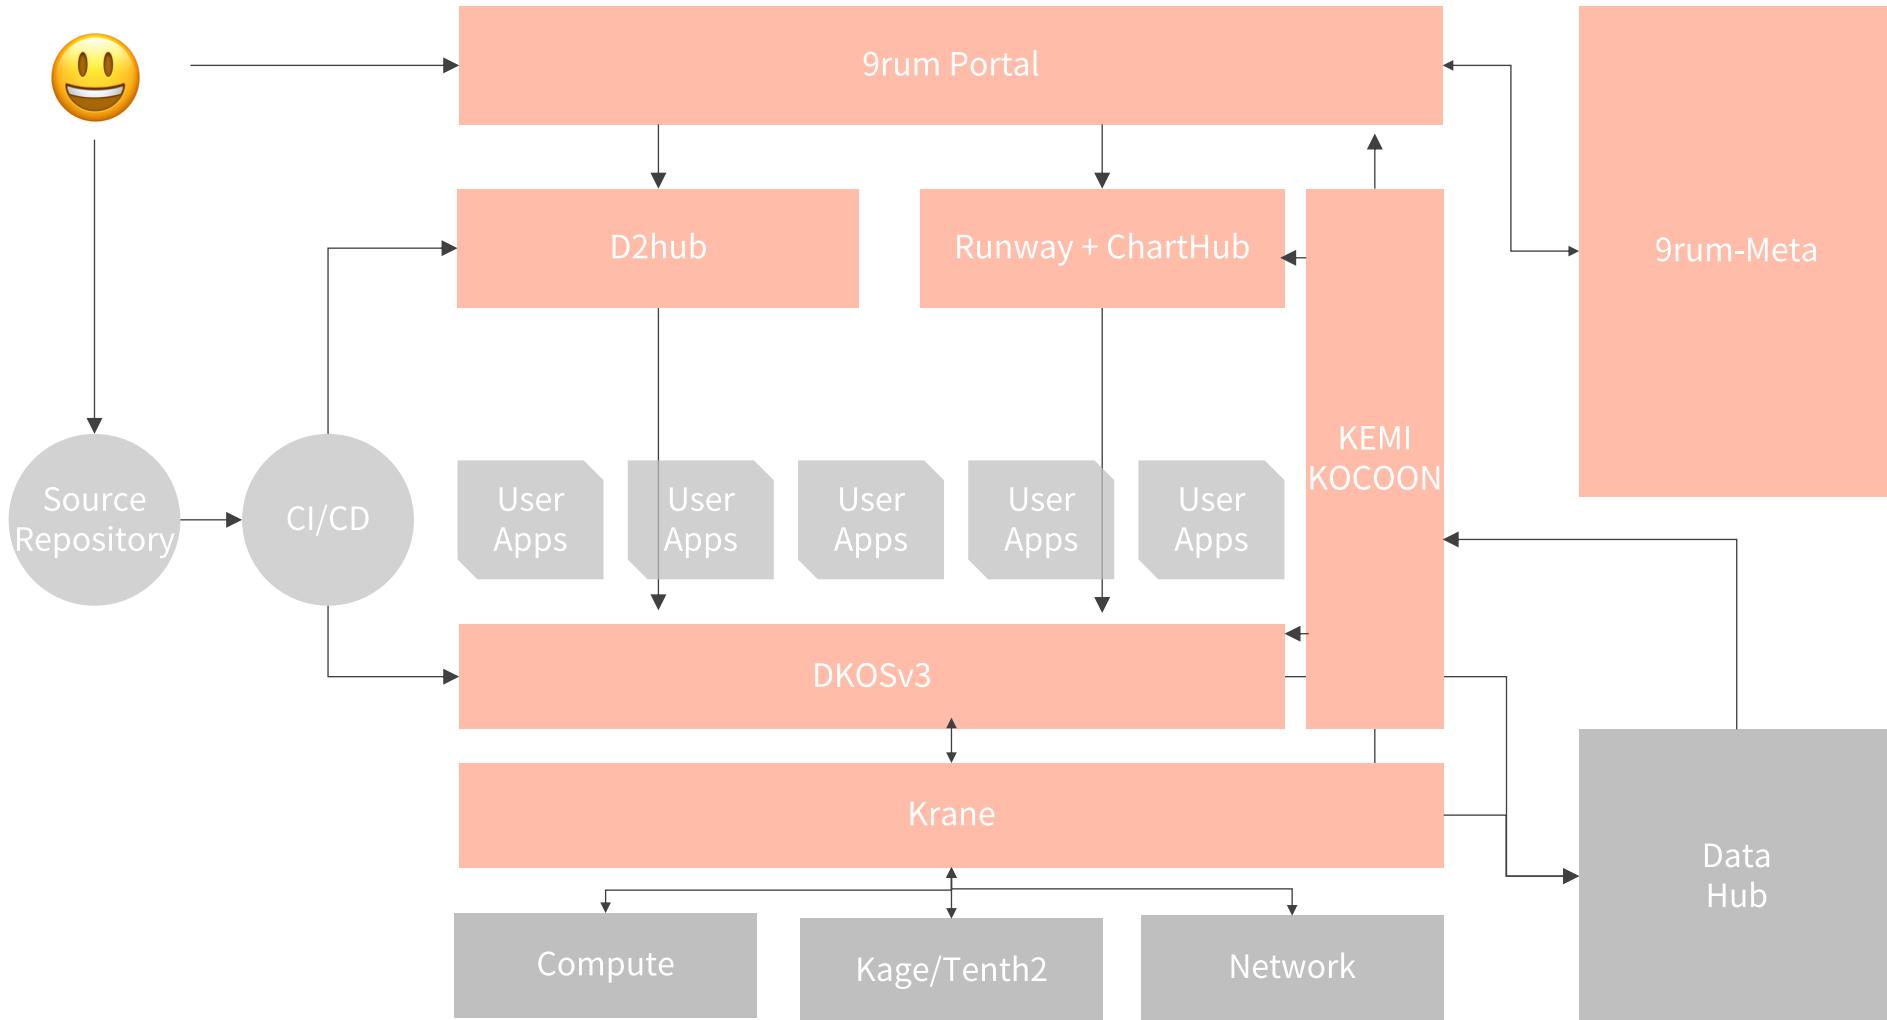


DKOSv3 Monitoring View



04 Summary

Again, Cloud Native Technologies at Kakao



Toward to Fully Automated Cloud

Journey might be hard

Cloud Native Technology에 대한
깊이 있는 이해 필요

Kubernetes 만으로는 부족, 다양한 리소스를
연결할 수 있어야 함

지속 가능한 기술을 확보하고 Cloud Native에 적합한
개발 문화로 바꾸기 위한 조직의 노력 필요

질문과 답변

감사합니다

we are hiring!